

Bike Renting Project

Madhurima Biswas

7 February 2019

Contents

I. Introduction	3
1.1 Problem Statement	3
1.2 Data	3
II. Methodology	5
2.1 Exploratory Data Analysis	5
2.1.1 Univariate Analysis	8
2.1.2 Bivariate Analysis	12
2.1.3 Data Consolidation	16
2.1.4 Missing Value Analysis	16
2.1.5 Outlier Analysis	16
2.1.6 Feature Selection	19
2.1.7 Feature Scaling	22
2.1.8 Data Sampling	22
2.2 Modeling	23
2.2.1 Model Development	23
2.2.2 Decision tree Algorithm	23
2.2.3 Random Forest Algorithm	26
2.2.4 Multiple Linear Regression	26
2.2.5 KNN Algorithm	29
III. Conclusion	30
3.1 Model Evaluation	30
3.2 Final Model Selection	31
 Appendix A - Extra Figures	 32
Appendix B - R Code	35
Appendix C - Python code	44
 References	 50

I. Introduction:

1.1. Problem Statement

Bike rental systems are a flexible transport service where users can rent a two-wheeler vehicle without going through the hassle of buying or maintaining one's own bike. We are provided daily rental data spanning two years 2011-2012. The objective of this case is to Predication of bike rental count on daily based on the environmental and seasonal settings, so that it helps in better management of the bike rental systems to organize & update their bikes for customers.

1.2. Data

The task is to build 'predictive regression' models, which will predict the bike rental count, based on the various factors given in the data during the year 2011-2012.

Given below is a sample of the data set that we are using to predict the rental count:

Table 1.1: Bike Renting sample data (year 2011-12)

instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit
1	01-01-11	1	0	1	0	6	0	2
2	02-01-11	1	0	1	0	0	0	2
3	03-01-11	1	0	1	0	1	1	1
4	04-01-11	1	0	1	0	2	1	1
5	05-01-11	1	0	1	0	3	1	1
6	06-01-11	1	0	1	0	4	1	1

temp	atemp	hum	windspeed	casual	registered	cnt
0.344167	0.363625	0.805833	0.160446	331	654	985
0.363478	0.353739	0.696087	0.248539	131	670	801
0.196364	0.189405	0.437273	0.248309	120	1229	1349
0.2	0.212122	0.590435	0.160296	108	1454	1562
0.226957	0.22927	0.436957	0.1869	82	1518	1600
0.204348	0.233209	0.518261	0.0895652	88	1518	1606

In the table above, we have the following 15 independent variables, using which we have to predict the Bike Rental Count:

Table 1.2: Predictor Variables

No. Independent Variables

1	instant
2	dteday
3	season
4	yr
5	mnth
6	holiday
7	weekday
8	workingday
9	weathersit
10	temp
11	atemp
12	hum
13	windspeed
14	casual
15	registered

We have in total 7 categorical variables, 8 numeric variables & one Date type variable.


II. Methodology

2.1. Exploratory Data Analysis:

The objective first is to study each feature available in the data and try to assess some patterns and understand the dimensions & properties of the data by exploring it visually. It helps us in understanding the nature of data in terms of distribution of the individual variables/features, finding missing values, relationship with other variables and many other things.

2.1.1. Univariate Analysis:

A. Dependent Target Variable: “cnt”



Since our target variable is continuous, we can visualize it by plotting its histogram.

Observation:

- The curve of the frequency distribution of “cnt” variable seems close to normal distribution curve, having mean = 4504 & median = 4548.
- Removing outliers might help reduce the slight skewness in data.
- Range: [22, 8714]

B. Independent Numeric Variables:

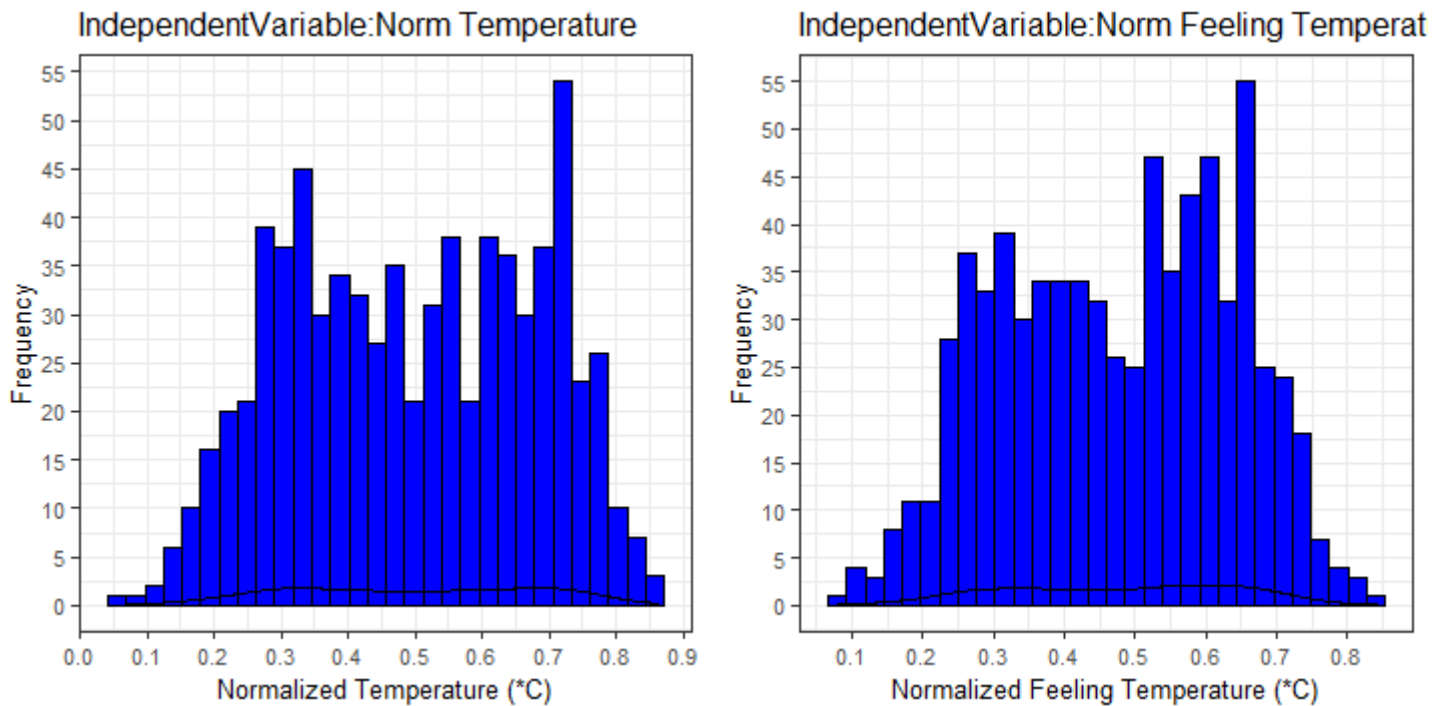


Fig.2.2

Observation:

- We see a wide range of temperature during years 2011-12; however, there is no clear-cut pattern.
- “Humidity” has a mean of 0.62 and “Windspeed” has a mean of 0.19 in 2011-12. Mostly are within a smaller range and the curve is slight skewed for both.


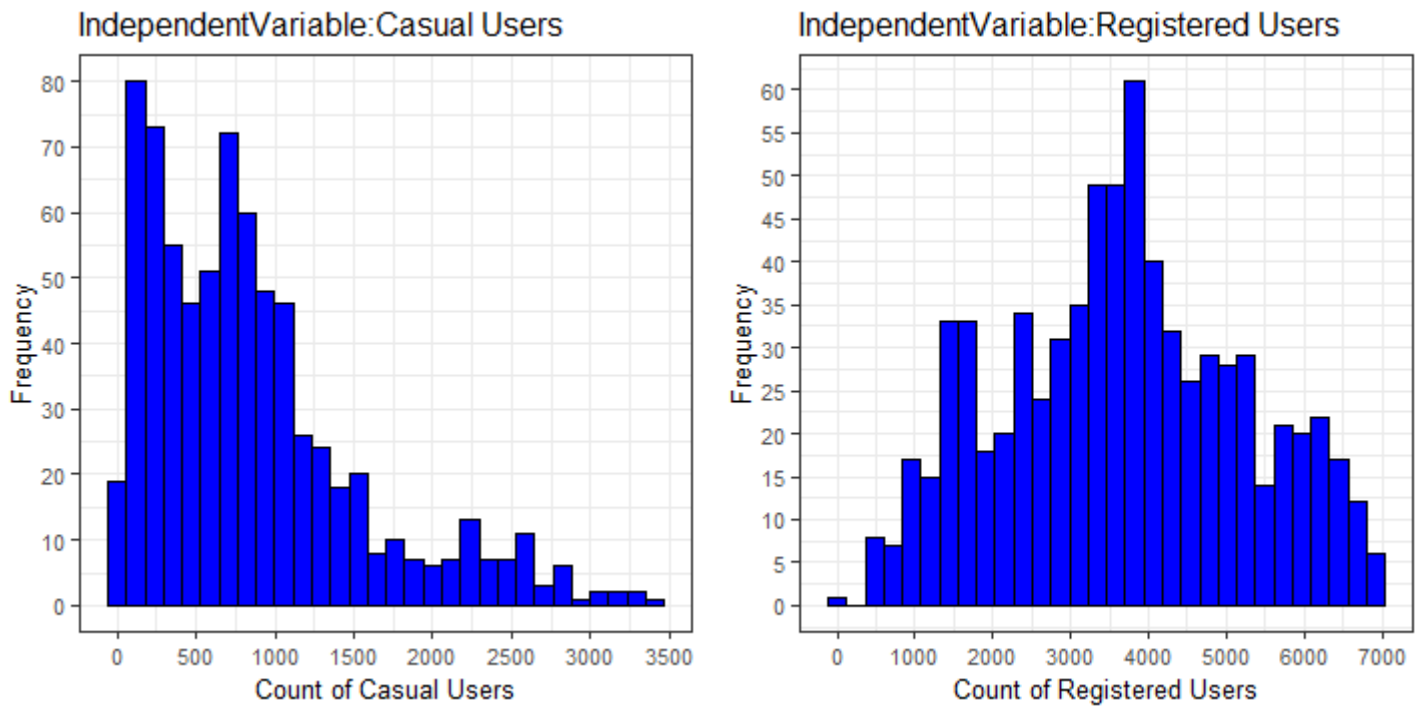


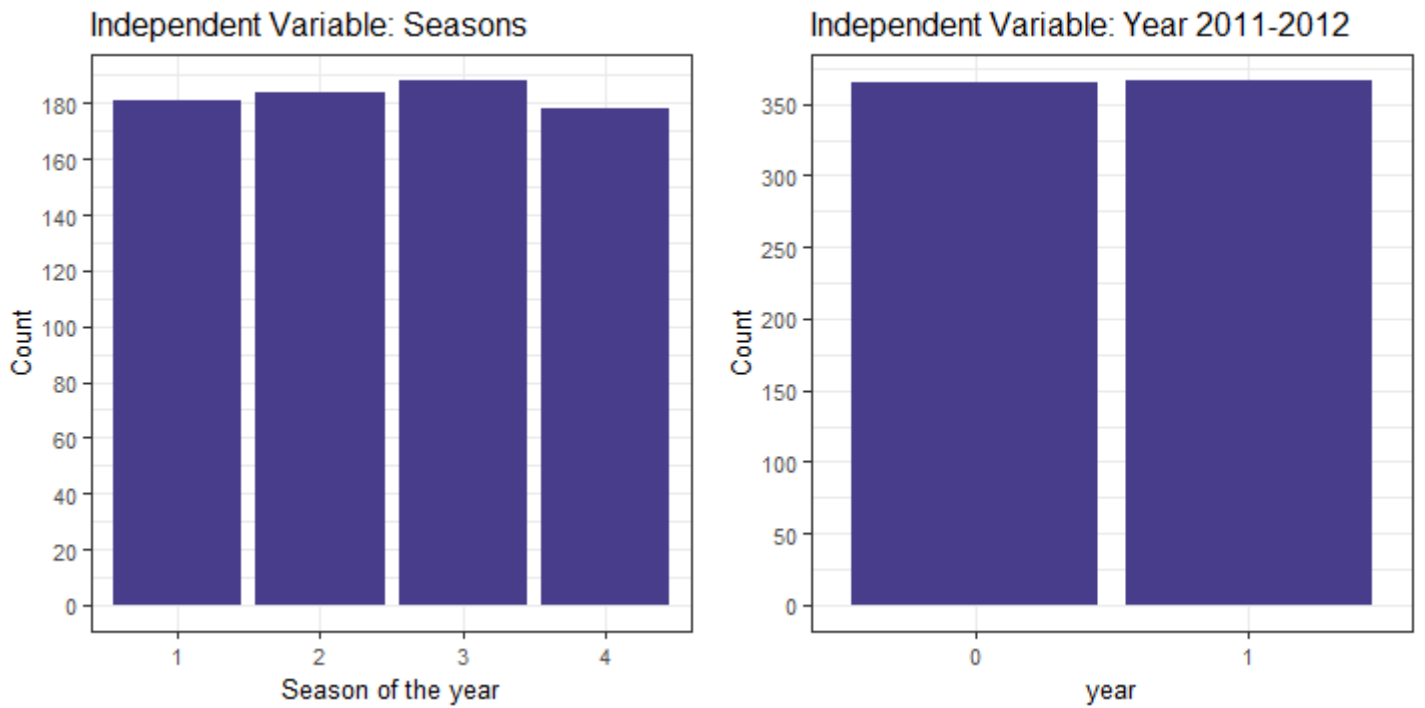
Fig.2.3



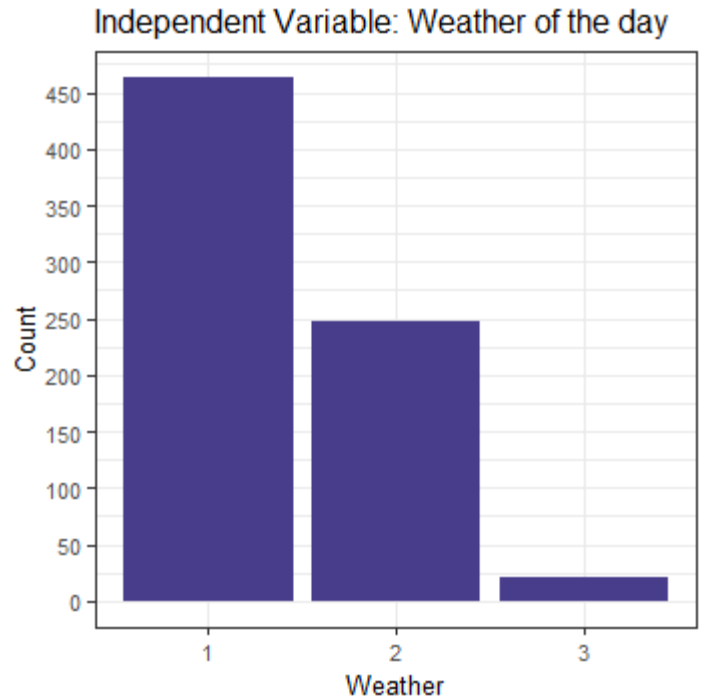
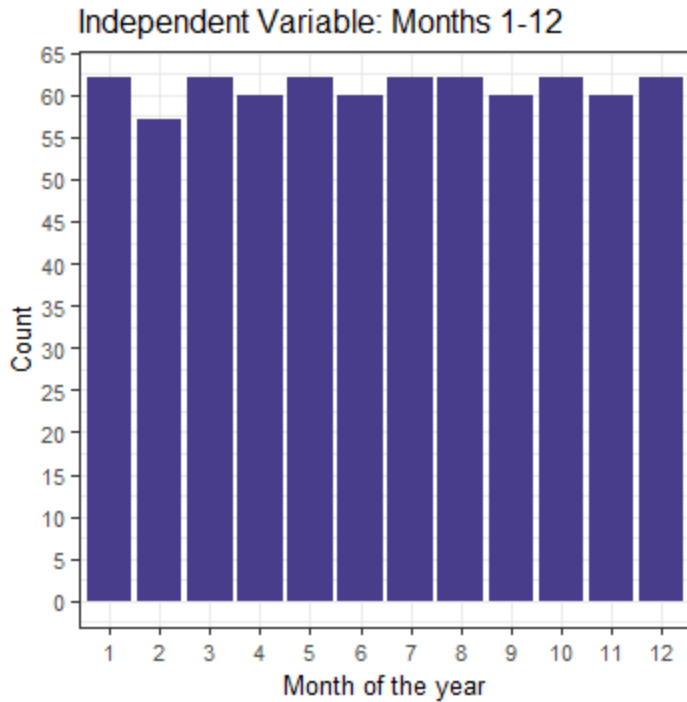
- We observe that most “Casual” per day are lesser in number and it is positively skewed frequency curve. “Registered” users per day have a mean of 3656 over 2011-12.

C. Independent Categorical Variables:

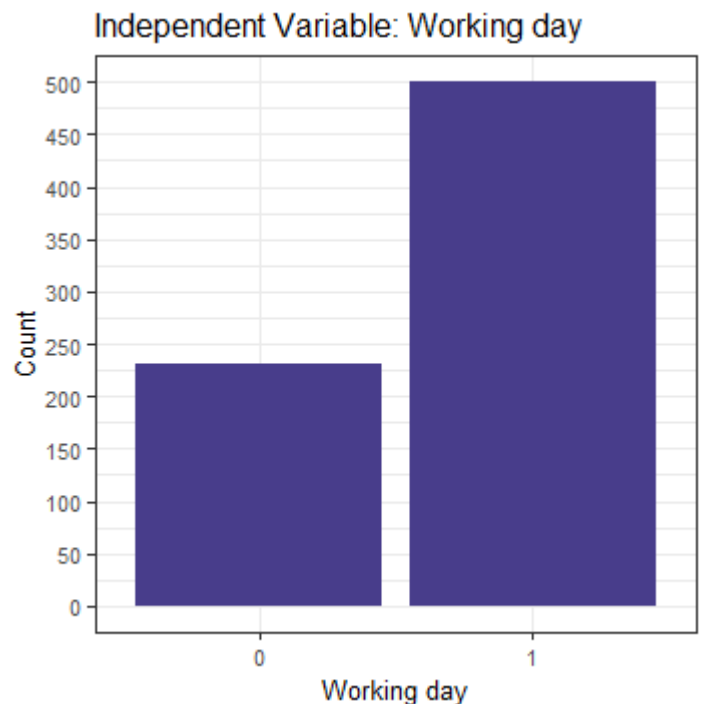
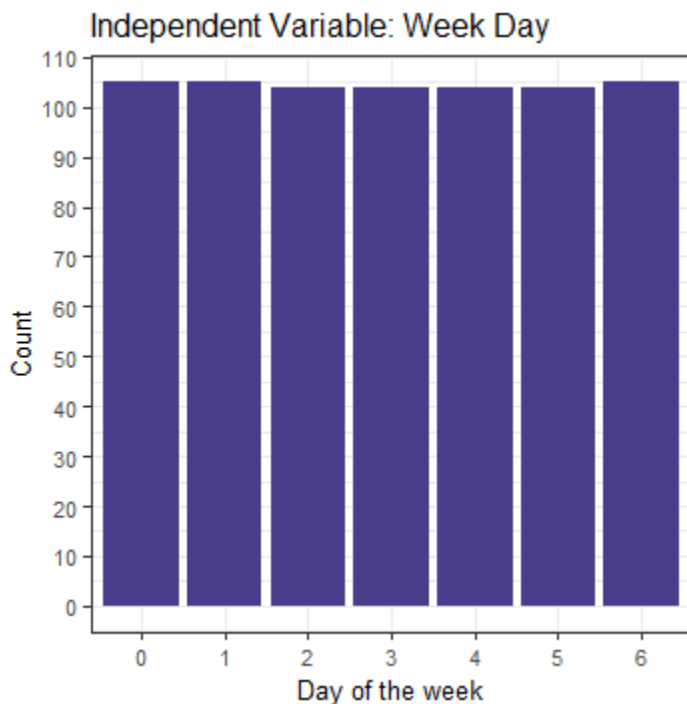
Bar graphs for the categorical variables in the data as follows:



- We observe that “fall” season was comparatively longer and “winter” was shortest in the years 2011-12.
- The data is well distributed within year 2011 & 2012 and we can hope to get less error due to data imbalance yearly.



- “Mnth” variable has correct data as February has least days and others have 30/31 days. The weather in 2011-12 was mostly clear or few clouds or partly cloudy.
- “week days” seem properly distributed and obviously there are more working days than holidays and weekends.




2.1.2. Bivariate Analysis:



After looking features individually, let's explore the independent variables with respect to the target variable using scatter plots to discover hidden relationships between the independent

variable and the target variable and use those findings in missing data imputation and feature engineering.


A. Dependent target Variable Vs Independent Variables:



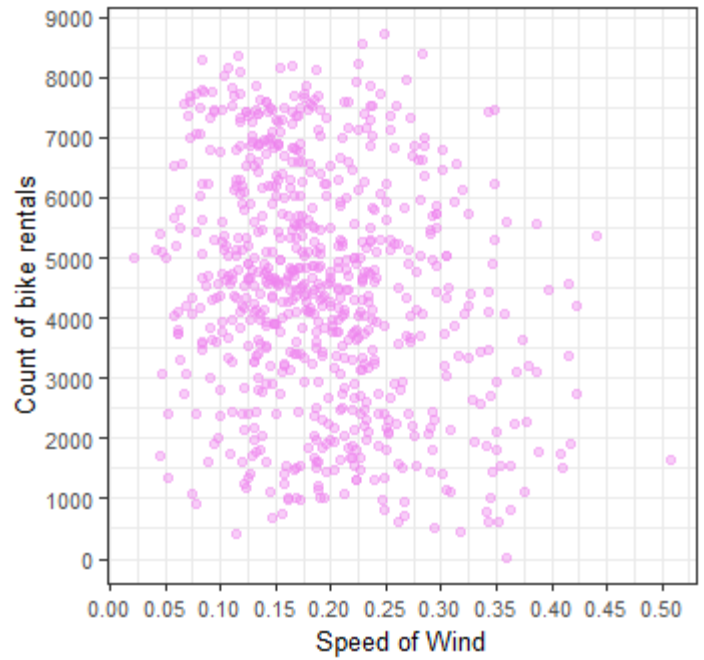
- We observe there are more bikes rented during the year 2012; it maybe because of the increase in popularity of the rental system after a successful year.
- There is no clear-cut pattern for temperature & rental count.



Scatter plot: Rental Count Vs Humidity

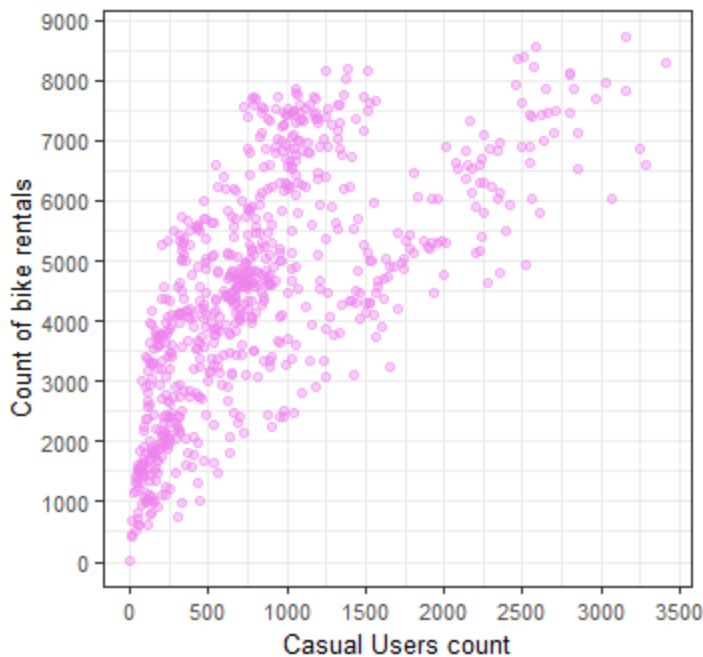


Scatter plot: Rental Count Vs Wind speed

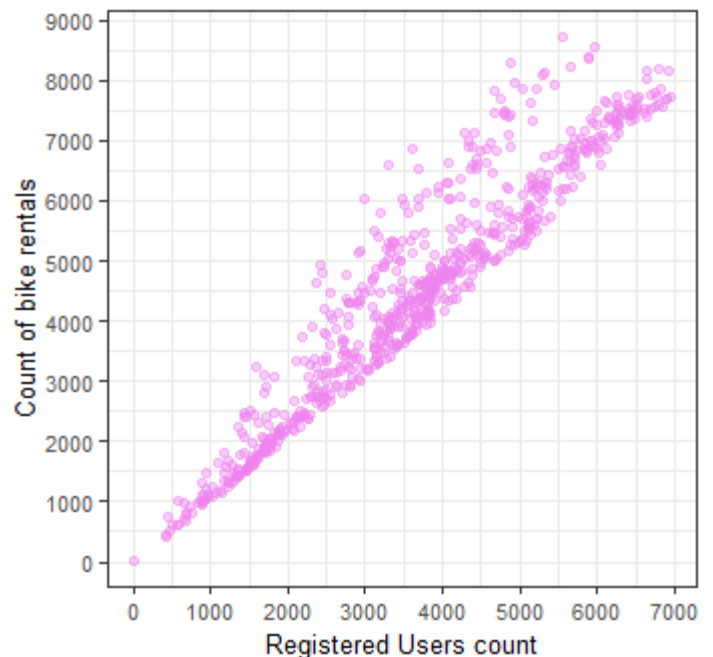


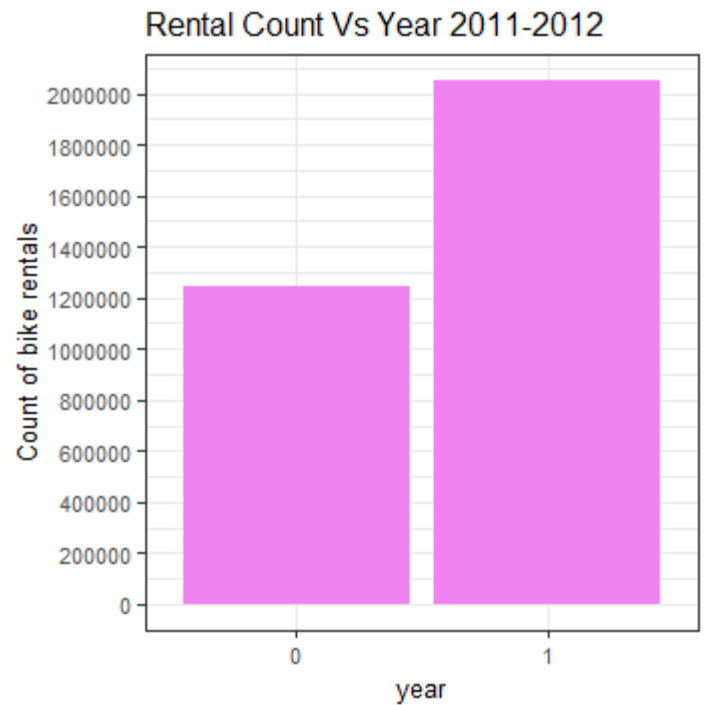
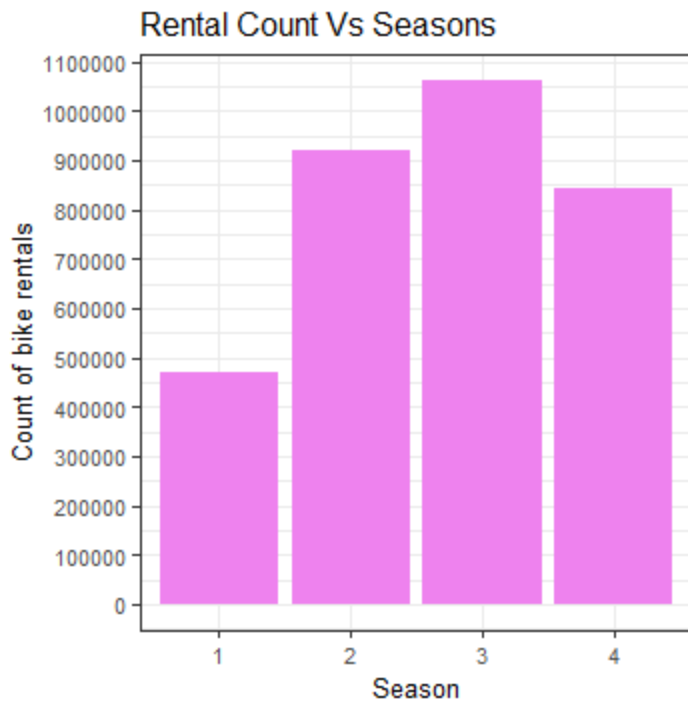
- No such clear-cut pattern for humidity or Wind speed & rental count.
- We see a slight linear relationship between Casual users & rental count and a strong linear relationship between Registered users & rental count, which is kind of obvious as rental count is the sum of total users, i.e. casual as well as registered customers.

Scatter plot: Rental Count Vs Casual users







Scatter plot: Rental Count Vs Registered users





- Rental count is most in fall season and that is expected as fall season was longest in 2011-12. However, the rental count is quite low in spring season which maybe due to low temperature during the season and there maybe seasonal influences on this.
- More bikes were rented in 2012 than in 2011, which could be due to the increase in popularity of bike rental systems.
- Most bikes rented in the month of June, July, August & September.

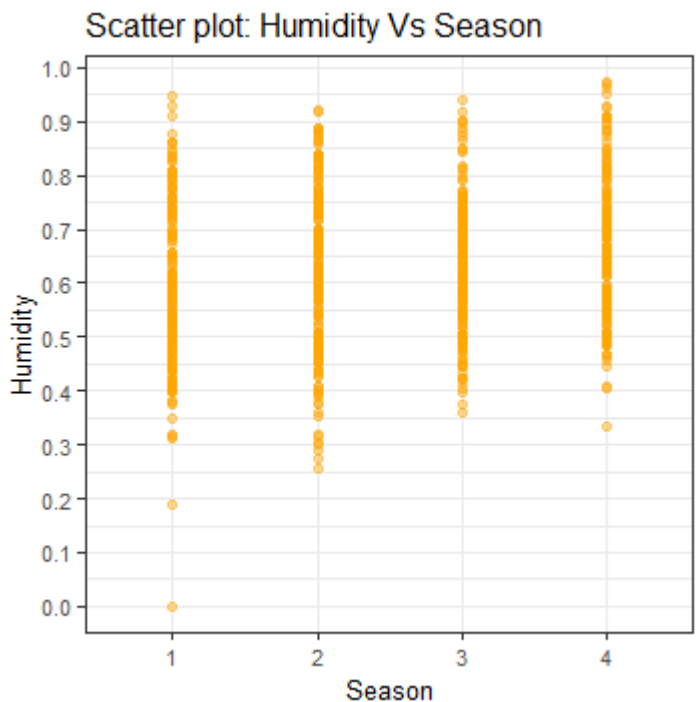
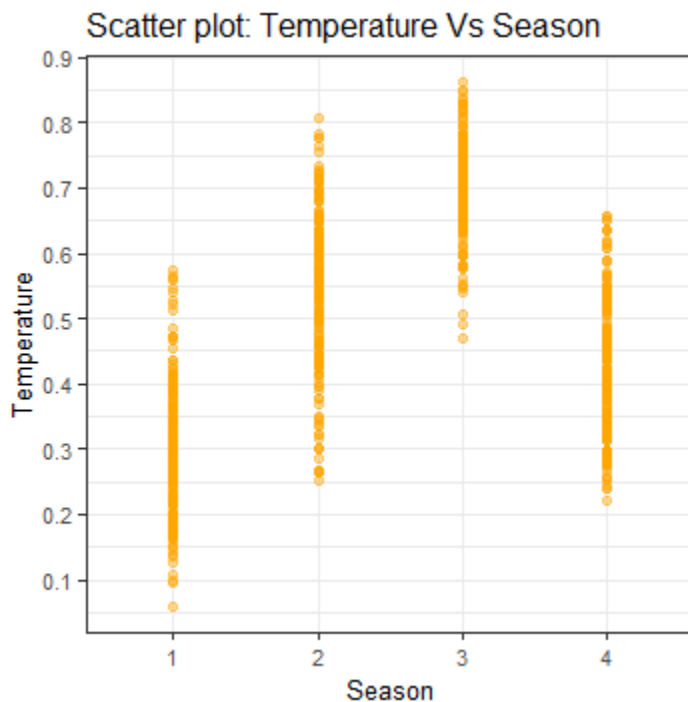




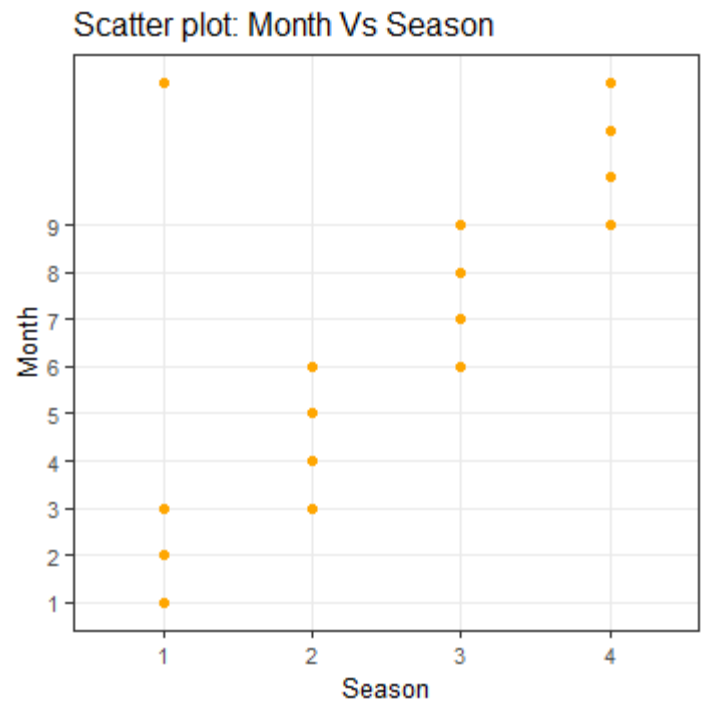
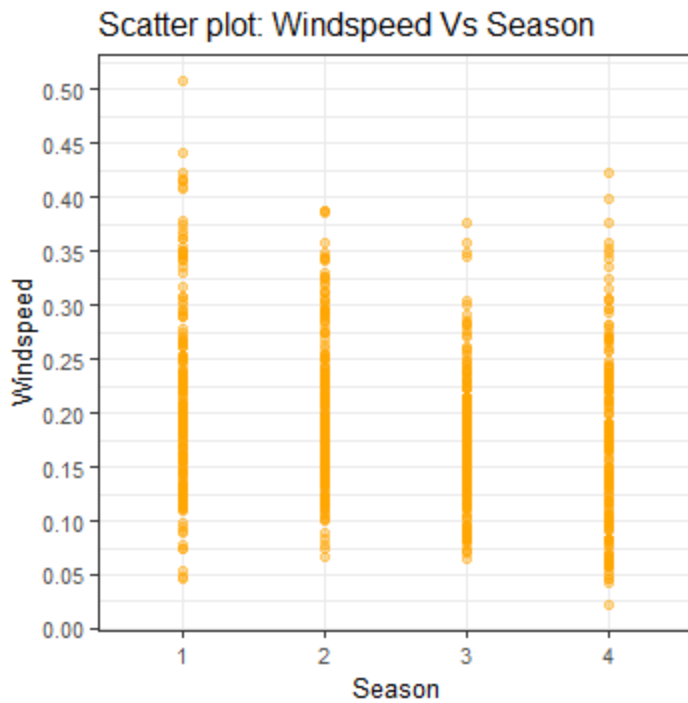
- Bikes are more rented on working days, but that is also because there are more working days in a year.
- More bikes rented on days with clear to partly cloudy weather, but that is also most days in a year.

B. Independent Variable Vs Independent Variable (Interdependencies):

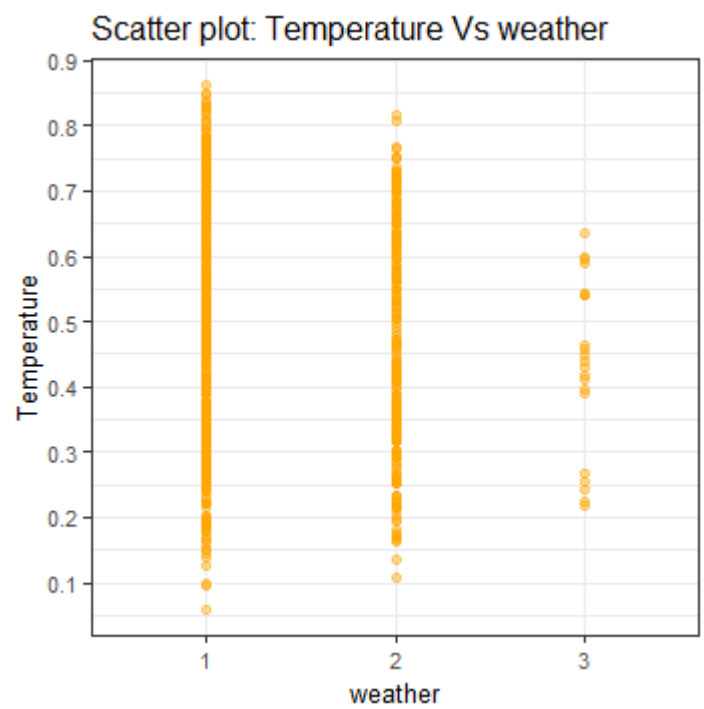
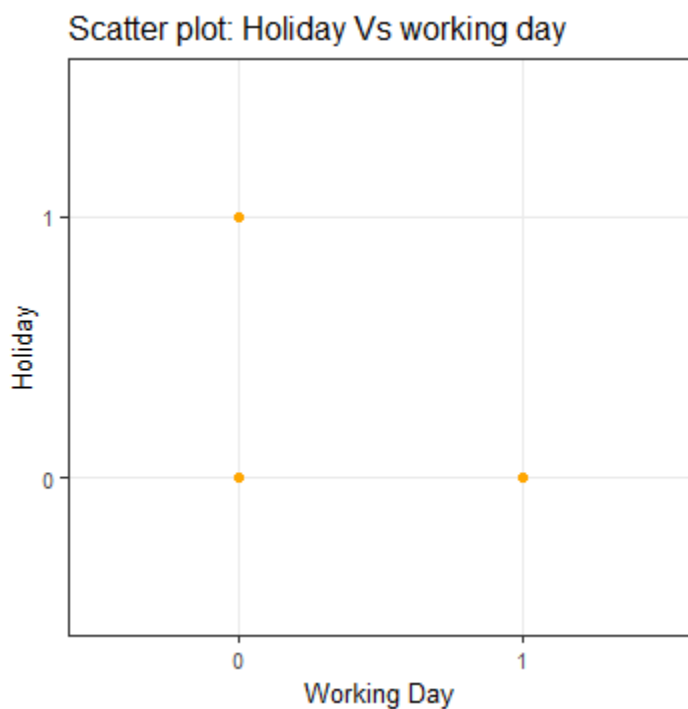
Let's explore the independent variables with respect to other independent variables using scatter plots to discover hidden relations or dependencies between them.





- We observe that temperature is highest in fall season and lowest in winter season.
- We don't see any dependence of humidity with seasonal changes.





- No such clear-cut pattern of windspeed due to seasonal changes.
- December, March, June & September are the months of seasonal transition.
- There are no days in the data where it was a holiday and a working day both simultaneously, as expected.
- During clear to partly cloudy days, the temperature has a wide range.






- There is a strong linear relationship between temperature & feeling temperature, as expected.
- No clear-cut pattern of humidity with temperature, except there were less dry air days in 2011-12.
- No clear-cut pattern of wind speed with temperature.
- There is a faint linear relationship between casual users & registered users, but not solid enough to be claimed.




Weather Vs Date



- Most of the year, we observed weather to be clear to partly cloudy in 2011-12.

Temperature Vs Date



- We observe a similar pattern of temperature in year 2011 & 2012, which tells us that the temperature changes were similar in both the year in changing seasons.

2.1.3. Data Consolidation:

Since not all data are in their proper data types, we need to convert it first to proceed further.

```
#_____Data type conversion_____#
catnames = c("season","yr","mnth","holiday","weekday","workingday","weathersit")
#categorical variables
for (i in catnames) {
  data[,i] = as.factor(data[,i])
}

numnames = c("temp","atemp","hum","windspeed","casual","registered","cnt")
#numerical variables
for (i in numnames) {
  data[,i] = as.numeric(data[,i])
}

data$dteday = as.Date(data$dteday)  #It changed date "02-04-11" to "2011-04-02".
```

2.1.4. Missing Value Analysis:

Missing data can have a severe impact on building predictive models because the missing values might contain some vital information, which could help in making better predictions. So, it becomes imperative to carry out missing data imputation.

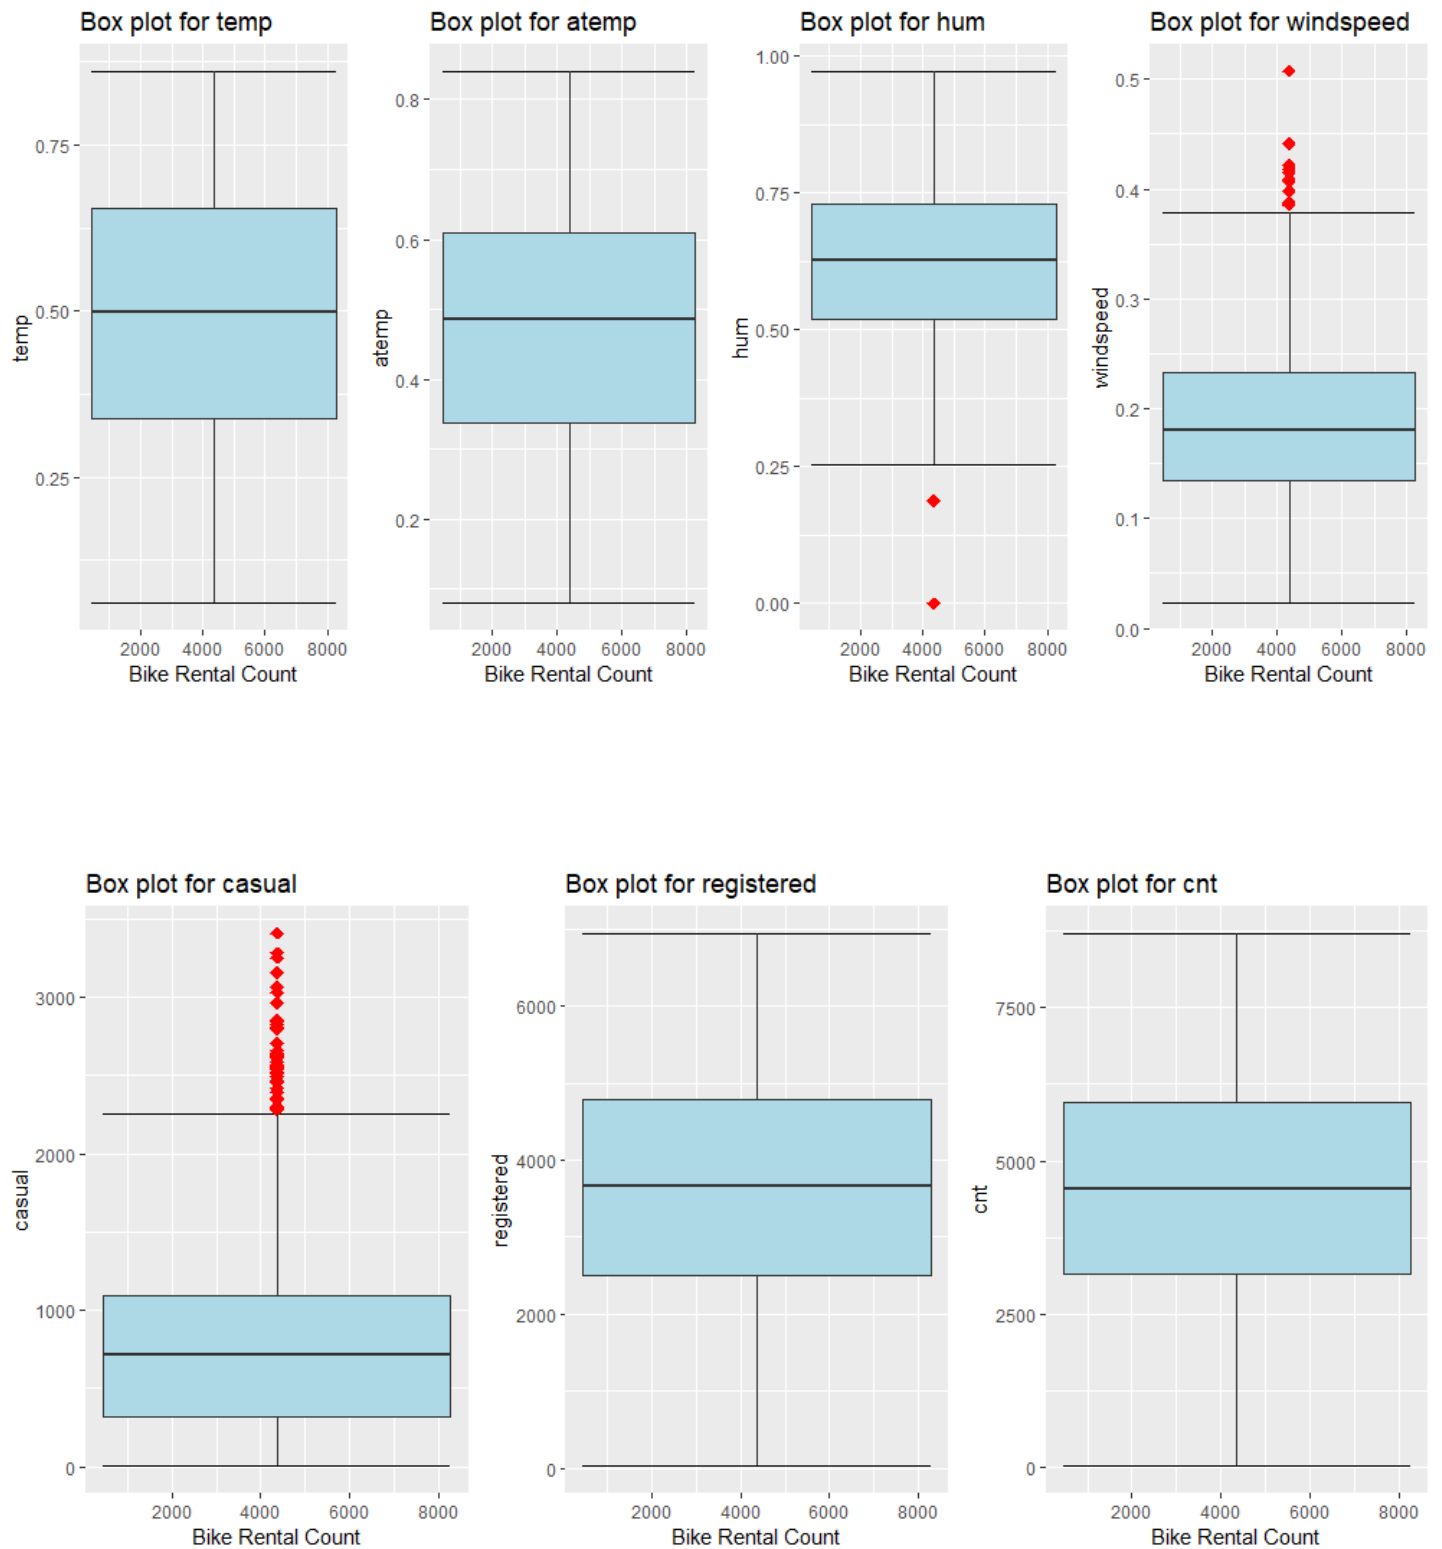
However, there are no missing values in this dataset and thus we move to next step.

2.1.5. Outlier Analysis:

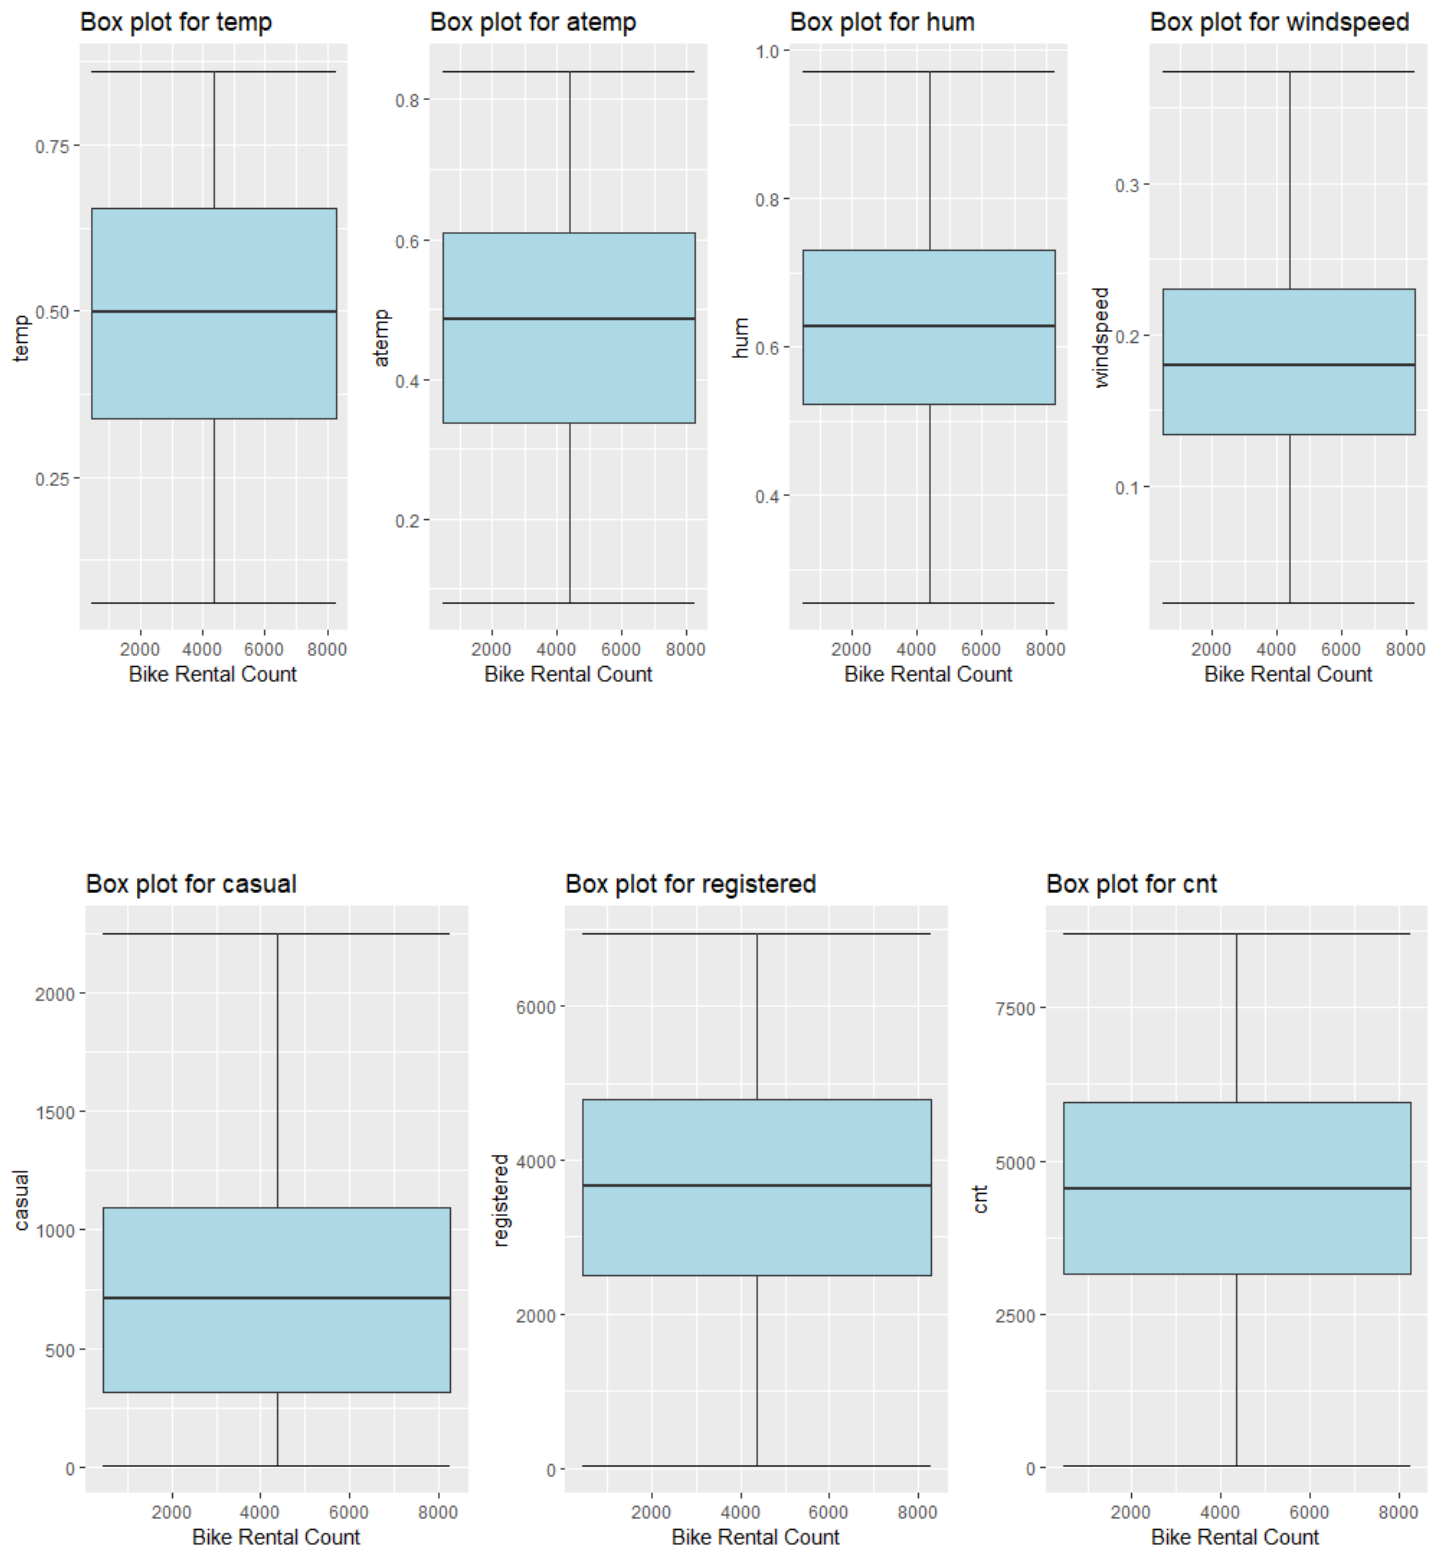
By definition, outliers are points that are distant from remaining observations. As a result, they can potentially skew or bias any analysis performed on the dataset. It is therefore important to detect and adequately deal with outliers using Box Plot method here.

Outliers make sense only in numeric or continuous data for this dataset. The “cnt” variable consists of the labels to be used to train and test the predictive models and hence it should be left untouched by further manipulation by outlier analysis.

Target variable “cnt” Box plots for Independent continuous Variables (before Outlier removal):




Box plots of the variables after Outlier removal:



2.1.6. Feature Selection:

It is needed that we assess the importance of each predictor variable in our analysis, as there is a possibility that many variables in our analysis are not important at all to predict the 'cnt' values.

A. Using Correlation plots:



- If $|r| > 0.8$ for two variables, those variables are considered redundant variables and one of them can be removed from the dataset.
- Output: "temp" & "atemp" variables are highly positively correlated as expected after performing the pre-processing of the data.
- Output: "cnt" & "registered" variables are highly positively correlated as expected after performing the pre-processing of the data.

B. Using Chi-square test of Independence (relationship between categorical variables):

(Dependencies amongst Independent Categorical variables)

#####Chi-square Test of Independence (within Categorical Variables)

```
for(i in catnames){  
  for(j in catnames){  
    if(i!=j){  
      print(names(data[i]))  
      print(paste0(" Vs ", names(data[j])))  
      print(chisq.test(table(data[,j],data[,i])))  
    }  
  }  
}
```

- If $p\text{-value} < 0.05$ (Reject Null Hypothesis) \Rightarrow Target variable depends on the independent variable.
- If $p\text{-value} > 0.05$ (Do Not Reject Null Hypothesis) \Rightarrow Target variable & independent variable are independent of each other.
- Output:

```
[1] "season"
```

```
[1] " Vs yr"
```

Pearson's Chi-squared test

```
data: table(data[, j], data[, i])
```

```
X-squared = 0.0041569, df = 3, p-value = 0.9999
```

```
[1] "season"
```

```
[1] " Vs mnth"
```

Pearson's Chi-squared test

```
data: table(data[, j], data[, i])
```

```
X-squared = 1765.1, df = 33, p-value < 2.2e-16
```

```
[1] "season"
```

```
[1] " Vs holiday"
```

Pearson's Chi-squared test

```
data: table(data[, j], data[, i])
```

```
X-squared = 1.4961, df = 3, p-value = 0.6832
```

```
.  
.  
.
```

- "workingday"-"holiday", "weekday"-"workingday", "weekday"-"holiday" & "mnth"-"season" depend on each other significantly.


C. Using Random Forest Algorithm:

#####Using Random Forest Algorithm:

```
data.rf=randomForest(data$cnt~.,data = data, ntree=1000, keep.forest= F, importance= T)
```

```
importance(data.rf,type = 1)
```

- "holiday" variable has the least importance.



[%IncMSE is the most robust and informative measure. It is the increase in mse of predictions (estimated with out-of-bag-CV) as a result of variable j being permuted (values randomly shuffled).]

D. Using ANOVA test (comparision of Target Vs categorical variables)

```
anovacat = aov(cnt ~ season + yr + mnth + holiday + workingday + weekday + weathersit , data = data)
```

```
summary(anovacat)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
season	3	950595868	316865289	436.234	< 2e-16 ***
yr	1	884008263	884008263	1217.030	< 2e-16 ***
mnth	11	187311622	17028329	23.443	< 2e-16 ***
holiday	1	3306975	3306975	4.553	0.03321 *
workingday	1	3209216	3209216	4.418	0.03591 *
weekday	5	12629845	2525969	3.478	0.00411 **
weathersit	2	185659616	92829808	127.800	< 2e-16 ***
Residuals	706	512813988	726365		

- If p-value<0.05 (Reject Null Hypothesis) => Population means are significantly different.
- If p-value>0.05 (Do Not Reject Null Hypothesis) => Population means are not significantly different or are same.

E. Feature Selection

We should remove those features that do not contribute to predicting the target variable as it will only lead to increase in the complexity of the model and reduce interpretability of models.

1. While doing data exploration, we notice that "instant" variable is just a serial number column, so we can remove it.
2. From Chi-square test, we notice that "working day", "holiday" & "weekday" depend on each other and intuitively there is a logical connection within them.

We make a new variable using this connection between the three variables

[Denote: 1-->weekend, 2--> working day, 3--> holiday]

```
data$day = NA
for (i in 1:nrow(data)){
  if ((data[i,7]=="0") && (data[i,5]=="0")){data[i,16] = 1}           #weekend
  else if ((data[i,7]=="1") && (data[i,5]=="0")){data[i,16] = 2}     #working day
  else if ((data[i,7]=="0") && (data[i,5]=="1")){data[i,16] = 3}     #holiday
  else data[i,16] =NA }
```

3. "Season" has multicollinearity problem as well and it is related to "mnth", so we can remove it.
4. "casual" & "registered" are basically the target variables as their addition results to "cnt". So, we can remove both & predict for just "cnt" variable.
5. "temp" & "atemp" are highly correlated and "atemp" variable's importance was found out to be more. Intuitively also, feeling temperature matters more for customers who will be travelling by bikes and hence "temp" variable is redundant.

```
data= subset(data, select= -c(season,workingday,temp,casual,registered))
```

After dimensional reduction, we have 731 observationa x 10 variables in our data set.

2.1.7. Feature Scaling:

The dataset contains features that are highly varying in magnitudes, units and range. Feature Scaling (Normalization/Standardization) is a step of Data Pre-Processing, which is applied to independent variables or features of data. It helps to normalize the data within a particular range and sometimes helps in speeding up the calculations in distance-based algorithms.

However, the continuous variables in the data set was already normalized.

2.1.8. Data Sampling:

The whole dataset is divided into train and test split sets so that there is data from which the model can learn and there is a part of the data set using which we can do unbiased evaluation of the trained model.

Random sampling without replacement is used to split 80% of the data into training set and remaining 20% into test set.

```
sample.index = sample(nrow(data), 0.8*nrow(data), replace = F) #80% data -->Train set, 20%-> Test set
train = data[sample.index,]
test = data[-sample.index,]
```

2.2. Modeling

2.2.1. Model Development:

The dataset of year 2011-2012 indicates that this is a supervised learning problem as there is the task of inferring a function or values from the labeled training data. Secondly, the dependent variable “cnt” is of real valued discrete type and therefore our prediction is of a quantity & it is a regression problem. Since we have many input variables, we shall perform a **multivariate regression analysis** on the given dataset.

2.2.2. Decision Tree Algorithm

Decision Trees

[Decision trees can handle both categorical and numerical variables at the same time as features. Every split in a decision tree is based on a feature. If the feature is categorical, the split is done with the elements belonging to a particular class. If the feature is continuous, the split is done with the elements higher than a threshold. At every split, the decision tree will take the best variable at that moment. This will be done according to an impurity measure with the splitted branches. And the fact that the variable used to do split is categorical or continuous is irrelevant (in fact, decision trees categorize continuous variables by creating binary regions with the threshold).]

```
dt=rpart(cnt~.,data = train,method= "anova")
> summary(dt)
Call:
rpart(formula = cnt ~ ., data = train, method = "anova")
n= 584

      CP nsplit rel error   xerror   xstd
1 0.37445616     0 1.0000000 1.0051616 0.04580505
2 0.22311915     1 0.6255438 0.6603832 0.03348656
3 0.09060873     2 0.4024247 0.4239814 0.03179904
4 0.02962425     3 0.3118160 0.3290237 0.02734505
5 0.02934392     4 0.2821917 0.3120647 0.02819117
6 0.02895436     5 0.2528478 0.3120647 0.02819117
7 0.01189898     6 0.2238934 0.2660670 0.02168208
8 0.01131214     7 0.2119945 0.2668795 0.02194647
9 0.01000000     8 0.2006823 0.2633306 0.02187781

Variable importance
      atemp      mnth      yr      hum  windspeed weathersit  weekday
      34       27      25       8         4         1         1

Node number 1: 584 observations,    complexity param=0.3744562
mean=4565.748, MSE=3745566
left son=2 (234 obs) right son=3 (350 obs)
Primary splits:
  atemp    < 0.4308565 to the left, improve=0.37445620, (0 missing)
  yr       splits as LR, improve=0.35623910, (0 missing)
  mnth     splits as LLLRRRRRRRLL, improve=0.30009300, (0 missing)
  weathersit splits as RLL, improve=0.07434951, (0 missing)
  hum      < 0.824394 to the right, improve=0.06695468, (0 missing)
Surrogate splits:
  mnth     splits as LLLRRRRRRRLL, agree=0.894, adj=0.735, (0 split)
  hum      < 0.5464585 to the left, agree=0.625, adj=0.064, (0 split)
  windspeed < 0.06282915 to the left, agree=0.616, adj=0.043, (0 split)
  dteday   < 29.5 to the right, agree=0.601, adj=0.004, (0 split)

Node number 2: 234 observations,    complexity param=0.09060873
mean=3117.359, MSE=2302852
left son=4 (126 obs) right son=5 (108 obs)
Primary splits:
```

```

yr          splits as LR, improve=0.36780560, (0 missing)
atemp       < 0.2607295 to the left, improve=0.23258030, (0 missing)
mnth        splits as LLLRL--R-RRR, improve=0.19311160, (0 missing)
hum         < 0.678777 to the right, improve=0.06662897, (0 missing)
weathersit   splits as RLL, improve=0.06151398, (0 missing)
Surrogate splits:
hum         < 0.5725 to the right, agree=0.577, adj=0.083, (0 split)
atemp       < 0.332973 to the left, agree=0.573, adj=0.074, (0 split)
windspeed   < 0.1871895 to the right, agree=0.568, adj=0.065, (0 split)
mnth        splits as LRLRL--R-LRL, agree=0.564, adj=0.056, (0 split)
weekday     splits as LLLLLRL, agree=0.543, adj=0.009, (0 split)

Node number 3: 350 observations, complexity param=0.2231192
mean=5534.1, MSE=2369868
left son=6 (164 obs) right son=7 (186 obs)
Primary splits:
yr          splits as LR, improve=0.58840310, (0 missing)
hum         < 0.834375 to the right, improve=0.15010660, (0 missing)
weathersit   splits as RRL, improve=0.09686697, (0 missing)
atemp       < 0.5018855 to the left, improve=0.06263038, (0 missing)
mnth        splits as -LRLRRRRRLR, improve=0.05588727, (0 missing)
Surrogate splits:
hum         < 0.6947915 to the right, agree=0.580, adj=0.104, (0 split)
mnth        splits as -RRLRLRRRRRLR, agree=0.569, adj=0.079, (0 split)
atemp       < 0.5296815 to the left, agree=0.549, adj=0.037, (0 split)
weekday     splits as RLLRRRR, agree=0.546, adj=0.030, (0 split)
windspeed   < 0.1741335 to the right, agree=0.543, adj=0.024, (0 split)

Node number 4: 126 observations, complexity param=0.02962425
mean=2265.302, MSE=1057926
left son=8 (75 obs) right son=9 (51 obs)
Primary splits:
mnth        splits as LLLLR----RRR, improve=0.48612910, (0 missing)
atemp       < 0.251738 to the left, improve=0.30669750, (0 missing)
windspeed   < 0.112571 to the right, improve=0.24712020, (0 missing)
hum         < 0.86 to the right, improve=0.11724950, (0 missing)
weathersit   splits as RLL, improve=0.07345125, (0 missing)
Surrogate splits:
windspeed   < 0.120031 to the right, agree=0.746, adj=0.373, (0 split)
atemp       < 0.298832 to the left, agree=0.714, adj=0.294, (0 split)
hum         < 0.611667 to the left, agree=0.611, adj=0.039, (0 split)
dteday      < 22.5 to the left, agree=0.603, adj=0.020, (0 split)
day         splits as LLR, agree=0.603, adj=0.020, (0 split)

Node number 5: 108 observations, complexity param=0.02895436
mean=4111.426, MSE=1920095
left son=10 (31 obs) right son=11 (77 obs)
Primary splits:
atemp       < 0.279985 to the left, improve=0.30542030, (0 missing)
mnth        splits as LLLR---R-LRL, improve=0.28345620, (0 missing)
hum         < 0.697292 to the right, improve=0.16823620, (0 missing)
weathersit   splits as RLL, improve=0.09756212, (0 missing)
weekday     splits as LLLRRRL, improve=0.07717721, (0 missing)
Surrogate splits:
hum         < 0.4647915 to the left, agree=0.741, adj=0.097, (0 split)
windspeed   < 0.349942 to the right, agree=0.731, adj=0.065, (0 split)
mnth        splits as RRRR---L-RRR, agree=0.722, adj=0.032, (0 split)
weathersit   splits as RRL, agree=0.722, adj=0.032, (0 split)

Node number 6: 164 observations, complexity param=0.01131214
mean=4276.524, MSE=648554.7
left son=12 (29 obs) right son=13 (135 obs)
Primary splits:
mnth        splits as -LLLRRRRRRLL, improve=0.23264010, (0 missing)

```



```

    hum          < 0.849375  to the right, improve=0.23168870, (0 missing)
    weathersit splits as RLL, improve=0.18122010, (0 missing)
    atemp        < 0.5805125  to the left,  improve=0.17080540, (0 missing)
    windspeed    < 0.1265645  to the right, improve=0.07228776, (0 missing)
  Surrogate splits:
    atemp        < 0.456723   to the left,  agree=0.872, adj=0.276, (0 split)
    windspeed    < 0.299444   to the right, agree=0.854, adj=0.172, (0 split)
    hum          < 0.908125   to the right, agree=0.829, adj=0.034, (0 split)

Node number 7: 186 observations,    complexity param=0.02934392
mean=6642.93, MSE=1263643
left son=14 (9 obs) right son=15 (177 obs)
Primary splits:
    hum          < 0.8322915  to the right, improve=0.27309330, (0 missing)
    weathersit splits as RLL, improve=0.13018900, (0 missing)
    atemp        < 0.4927355  to the left,  improve=0.12328470, (0 missing)
    mnth         splits as -LLLRRRRRR-L, improve=0.07749548, (0 missing)
    windspeed    < 0.287627   to the right, improve=0.06415826, (0 missing)
  Surrogate splits:
    weathersit splits as RRL, agree=0.968, adj=0.333, (0 split)
    windspeed    < 0.3526145  to the right, agree=0.957, adj=0.111, (0 split)

Node number 8: 75 observations
mean=1673.933, MSE=304991.8

Node number 9: 51 observations
mean=3134.961, MSE=894587.3

Node number 10: 31 observations
mean=2904.516, MSE=1394240

Node number 11: 77 observations,    complexity param=0.01189898
mean=4597.325, MSE=1309269
left son=22 (18 obs) right son=23 (59 obs)
Primary splits:
    hum          < 0.700625   to the right, improve=0.25817860, (0 missing)
    mnth         splits as LLLR-----LRL, improve=0.23626120, (0 missing)
    weathersit splits as RL-, improve=0.15559330, (0 missing)
    atemp        < 0.3134065  to the left,  improve=0.08333220, (0 missing)
    dteday       < 19.5       to the right, improve=0.07422147, (0 missing)
  Surrogate splits:
    weathersit splits as RL-,          agree=0.792, adj=0.111, (0 split)
    mnth         splits as RRRR-----LRR, agree=0.779, adj=0.056, (0 split)

Node number 12: 29 observations
mean=3438.448, MSE=473523.1

Node number 13: 135 observations
mean=4456.556, MSE=502863

Node number 14: 9 observations
mean=4037.778, MSE=2317994

Node number 15: 177 observations
mean=6775.395, MSE=847392.4

Node number 22: 18 observations
mean=3544.722, MSE=1303907

Node number 23: 59 observations
mean=4918.458, MSE=869753.4

```

We predict for test set:

```
predict.dt=predict(dt,test[, -10])
```

2.2.3. Random Forest Algorithm

[Random forests are an ensemble learning method for classification, regression and other tasks that operates by constructing a multitude of decision trees at training time and outputting the class that is the mode of the classes (classification) or mean prediction (regression) of the individual trees.]

```
rf = randomForest(cnt~., train, importance = TRUE, ntree = 500)
> summary(rf)
```

	Length	Class	Mode
call	5	-none-	call
type	1	-none-	character
predicted	584	-none-	numeric
mse	500	-none-	numeric
rsq	500	-none-	numeric
oob.times	584	-none-	numeric
importance	18	-none-	numeric
importanceSD	9	-none-	numeric
localImportance	0	-none-	NULL
proximity	0	-none-	NULL
ntree	1	-none-	numeric
mtry	1	-none-	numeric
forest	11	-none-	list
coefs	0	-none-	NULL
y	584	-none-	numeric
test	0	-none-	NULL
inbag	0	-none-	NULL
terms	3	terms	call

We predict for test set:

```
predict.rf <- data.frame(predict(rf, subset(test, select = -c(cnt))))
```

2.2.4. Multiple Linear Regression

Multicollinearity is when independent variables in a regression model are correlated. It tries to inflate or resist the variance of different strong regressors in the data. Therefore, we need to do a collinearity check before performing linear regression.

```
#creating dummy variables for categorical data
factor_new = dummy.data.frame(factor_data, sep = ".") #731 x 27
>
> #sampling#
> df = cbind(factor_new, num_data)
> #for (i in 1:ncol(df)) {
> #  df[,i] = as.numeric(df[,i])
> #}
> str(df) # 731 x 32
'data.frame': 731 obs. of 32 variables:
 $ yr.0 : int 1 1 1 1 1 1 1 1 1 1 ...
 $ yr.1 : int 0 0 0 0 0 0 0 0 0 0 ...
 $ mnth.1 : int 1 1 1 1 1 1 1 1 1 1 ...
 $ mnth.2 : int 0 0 0 0 0 0 0 0 0 0 ...
 $ mnth.3 : int 0 0 0 0 0 0 0 0 0 0 ...
 $ mnth.4 : int 0 0 0 0 0 0 0 0 0 0 ...
 $ mnth.5 : int 0 0 0 0 0 0 0 0 0 0 ...
 $ mnth.6 : int 0 0 0 0 0 0 0 0 0 0 ...
 $ mnth.7 : int 0 0 0 0 0 0 0 0 0 0 ...
 $ mnth.8 : int 0 0 0 0 0 0 0 0 0 0 ...
```

```

$ mnth.9      : int  0 0 0 0 0 0 0 0 0 0 ...
$ mnth.10     : int  0 0 0 0 0 0 0 0 0 0 ...
$ mnth.11     : int  0 0 0 0 0 0 0 0 0 0 ...
$ mnth.12     : int  0 0 0 0 0 0 0 0 0 0 ...
$ day.1       : int  1 1 0 0 0 0 0 1 1 0 ...
$ day.2       : int  0 0 1 1 1 1 1 0 0 1 ...
$ day.3       : int  0 0 0 0 0 0 0 0 0 0 ...
$ weekday.0   : int  0 1 0 0 0 0 0 0 1 0 ...
$ weekday.1   : int  0 0 1 0 0 0 0 0 0 1 ...
$ weekday.2   : int  0 0 0 1 0 0 0 0 0 0 ...
$ weekday.3   : int  0 0 0 0 1 0 0 0 0 0 ...
$ weekday.4   : int  0 0 0 0 0 1 0 0 0 0 ...
$ weekday.5   : int  0 0 0 0 0 0 1 0 0 0 ...
$ weekday.6   : int  1 0 0 0 0 0 0 1 0 0 ...
$ weathersit.1: int  0 0 1 1 1 1 0 0 1 1 ...
$ weathersit.2: int  1 1 0 0 0 0 1 1 0 0 ...
$ weathersit.3: int  0 0 0 0 0 0 0 0 0 0 ...
$ dteday      : num  1 2 3 4 5 6 7 8 9 10 ...
$ atemp       : num  0.364 0.354 0.189 0.212 0.229 ...
$ hum         : num  0.806 0.696 0.437 0.59 0.437 ...
$ windspeed   : num  0.16 0.249 0.248 0.16 0.187 ...
$ cnt         : num  985 801 1349 1562 1600 ...
>
> set.seed(123)
> train_index = sample(1:nrow(df), 0.8*nrow(df))
> train.df = df[train_index,]          #584 x 32
> test.df = df[-train_index,]         #147 x 32
>
> #Check Multicollinearity
vif(df[, -32])
  Variables      VIF
1      yr.0      Inf
2      yr.1      Inf
3    mnth.1      Inf
4    mnth.2      Inf
5    mnth.3      Inf
6    mnth.4      Inf
7    mnth.5      Inf
8    mnth.6      Inf
9    mnth.7      Inf
10   mnth.8      Inf
11   mnth.9      Inf
12  mnth.10      Inf
13  mnth.11      Inf
14  mnth.12      Inf
15    day.1      Inf
16    day.2      Inf
17    day.3      Inf
18 weekday.0      Inf
19 weekday.1      Inf
20 weekday.2      Inf
21 weekday.3      Inf
22 weekday.4      Inf
23 weekday.5      Inf
24 weekday.6      Inf
25 weathersit.1      Inf
26 weathersit.2      Inf
27 weathersit.3      Inf
28      dteday 1.010204
29      atemp 6.049203
30      hum 2.294781
31  windspeed 1.207595
> vifcor(df[, -32], th = 0.8)
3 variables from the 31 input variables have collinearity problem:

```

yr.1 weathersit.2 day.2

After excluding the collinear variables, the linear correlation coefficients ranges between:

min correlation (windspeed ~ weekday.3): -0.0001206042

max correlation (weekday.0 ~ day.1): 0.6450846

----- VIFs of the remained variables -----

	Variables	VIF
1	yr.0	1.049547
2	mnth.1	Inf
3	mnth.2	Inf
4	mnth.3	Inf
5	mnth.4	Inf
6	mnth.5	Inf
7	mnth.6	Inf
8	mnth.7	Inf
9	mnth.8	Inf
10	mnth.9	Inf
11	mnth.10	Inf
12	mnth.11	Inf
13	mnth.12	Inf
14	day.1	Inf
15	day.3	1.106961
16	weekday.0	Inf
17	weekday.1	Inf
18	weekday.2	Inf
19	weekday.3	Inf
20	weekday.4	Inf
21	weekday.5	Inf
22	weekday.6	Inf
23	weathersit.1	1.779943
24	weathersit.3	1.222714
25	dteday	1.010204
26	atemp	6.049203
27	hum	2.294781
28	windspeed	1.207595

> #Output:

> #3 variables from the 31 input variables have collinearity problem: yr.1, weathersit.2, day.2

> #removing multicollinear variables and redo check:

> df = subset(df, select= -c(yr.1, weathersit.2, day.2))

> train.df = subset(train.df, select= -c(yr.1, weathersit.2, day.2)) #584 x 29

> test.df = subset(test.df, select= -c(yr.1, weathersit.2, day.2)) #147 x 29

> dim(df) #731 x 29

[1] 731 29

> #Recheck VIFCORR: No variable from the 29 input variables has collinearity problem.

>

> #run regression model

> lr = lm(cnt~., data = train.df)

> #summary of the model

> summary(lr)

Call:

lm(formula = cnt ~ ., data = train.df)

Residuals:

Min	1Q	Median	3Q	Max
-3876.2	-387.8	50.8	509.4	2771.2

Coefficients: (3 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4430.566	344.461	12.862	< 2e-16 ***

```

yr.0      -2113.166      71.121 -29.712 < 2e-16 ***
mnth.1    -825.824      175.718 -4.700 3.29e-06 ***
mnth.2    -716.510      179.767 -3.986 7.62e-05 ***
mnth.3      138.034      174.608  0.791 0.429552
mnth.4      632.261      191.820  3.296 0.001043 **
mnth.5      957.277      209.921  4.560 6.29e-06 ***
mnth.6      673.222      240.603  2.798 0.005319 **
mnth.7      362.334      258.956  1.399 0.162305
mnth.8      644.409      241.596  2.667 0.007868 **
mnth.9     1396.680      213.404  6.545 1.35e-10 ***
mnth.10    1391.067      187.618  7.414 4.56e-13 ***
mnth.11     785.587      172.682  4.549 6.61e-06 ***
mnth.12      NA          NA      NA      NA
day.1       8.810      129.991  0.068 0.945990
day.3     -813.416      212.812 -3.822 0.000147 ***
weekday.0  -424.315      129.802 -3.269 0.001146 **
weekday.1  -165.593      133.805 -1.238 0.216395
weekday.2  -151.960      130.711 -1.163 0.245504
weekday.3   -23.876      130.518 -0.183 0.854920
weekday.4   -54.480      133.389 -0.408 0.683114
weekday.5      NA          NA      NA      NA
weekday.6      NA          NA      NA      NA
weathersit.1  448.480      95.588  4.692 3.41e-06 ***
weathersit.3 -1468.420     232.217 -6.323 5.24e-10 ***
dteday     -10.119        3.989 -2.537 0.011455 *
atemp      4592.019      519.614  8.837 < 2e-16 ***
hum        -1522.767      365.632 -4.165 3.61e-05 ***
windspeed   -2629.300      543.404 -4.839 1.69e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 837.3 on 558 degrees of freedom
Multiple R-squared: 0.8237, Adjusted R-squared: 0.8158
F-statistic: 104.3 on 25 and 558 DF, p-value: < 2.2e-16

We predict for test set:

```
predict.lm = predict(lm, test.df[, -29])
```

2.2.5. KNN Implementation


KNN is distance based non-parametric algorithm and it never stores patterns from the training data, but classifies for new test cases based on a similarity measure.

First, we need to check for the best no. of neighbors (k):

```

#To check for best k value:
model <- train(cnt~, data = train, method = "knn",
  trControl = trainControl("cv", number = 10),
  tuneLength = 15)
model$bestTune
#k = 3, 9
plot(model)

```



After checking both methods, it is best to choose $k=3$ as it gives us the least prediction error.

III. Conclusion

3.1 Model Evaluation:

Now that we have a few models for predicting the target variable, we need to decide which one to choose. Several criteria exist for evaluating and comparing models; here we can compare the models by using assessing the 'Predictive Performance' of the models. Predictive performance can be measured by comparing Predictions of the models with real values of the target variables, and calculating some average error measure like RMSE or MAPE.

#Error metric for Decision Tree:

```
postResample(predict.dt,test[,10])
```

#Output:

```
#RMSE    Rsquared    MAE
```

```
#1036.8218286  0.7105788 768.8217306
```

#Error metric for Random Forest:

```
postResample(predict.rf,test[,10])
```

#Output:

```
#RMSE    Rsquared    MAE
```

```
#778.4675527  0.8507608 576.6110231
```

#Error metric for Multiple Linear Regression:

```
postResample(predict.lr,test.df[,29])
```

#Output:

```
#RMSE    Rsquared    MAE
```


```
#800.2783046  0.8303233 581.4298996
```

```
#Error metric for KNN:
postResample(predict.knn$pred,test.df[,29])
#Output:
#RMSE      Rsquared      MAE
#1392.7631351  0.4544424 1110.0045351

#calculate MAPE
> mape = function(y,yi)
+ {mean(abs((y-yi)/y))*100
+ }
> mape.dt = mape(test[,10],predict.dt) #30.79%
> mape.rf = mape(test[,10],predict.rf$predict.rf..subset.test..select....c.cnt...) # 24.9%
> mape.lr = mape(test.df[,29],predict.lr) #17.5%
> mape.knn = mape(test.df[,29],predict.knn$pred) #38.98%
```




```
>      algorithm MAPE_val
1      Decision Tree 30.79662
2      Random Forest 24.98612
3 Linear Regression 17.55068
4              KNN 38.98097
```



3.2 Final Model Selection:




As we can observe that “Multiple Linear Regression” algorithm produces the least error or MAPE (Mean Absolute Percentage Error), we can freeze this algorithm as the model for analysis of new daily data or test cases of Bike Rental count for further years.

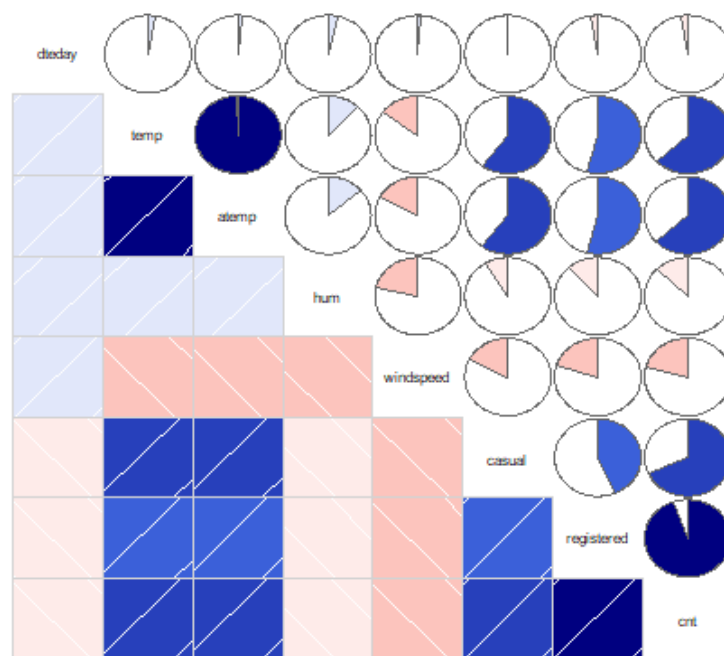
Appendix A: Extra plots







Correlation Plot



Appendix B: R code

```
#To clear the R environment of any predefined objects
rm(list=ls())

#To set working directory
setwd("F:/DS/edWisor/Project 2")
getwd()

#To load required libraries
library(ggplot2) # used for plotting
library(dplyr)   # used for data manipulation and joining
library(scales)  # used for "pretty_brakes()" function"
library(DMwR)    # used for KNN Imputation
library(outliers) # used for outlier detection & modification
library(corrgram) # used for plotting correlation amongst variables
library(corrplot) # used for plotting correlation amongst variables
library(caret)   # used for various model training
library(lubridate) # used for handling date format data
library(FNN)     # used for KNN modeling
library(randomForest) # used for Random Forest implementation
library(rpart)   # used for Decision Tree algorithm implementation

#To load the data
data = read.csv("day.csv",header = T, na.strings = c("", " ", "NA", NA))

#####Data Exploration#####
str(data)      #"data.frame"
dim(data)      # 731 x 16

###Univariate Analysis###
#col = names(data)
#To find the unique values in each column
#for (i in col) {
# print(i)
# print(length(unique(data[,i])))
#}

#Data has 7 categorical variables, 8 numeric variables & one date type variable.
#Target variable is integer type in nature.

###Data Consolidation###
#Convert into Proper data types
#-->ignoring "instant" as it is just like serial number.
data = data[,-1]
#dim(data)      #731 x 15
```

```

#_____Data type conversion_____#
catnames = c("season","yr","mnth","holiday","weekday","workingday","weathersit") #categorical variables
for (i in catnames) {
  data[,i] = as.factor(data[,i])
}

numnames = c("temp","atemp","hum","windspeed","casual","registered","cnt")      #numerical variables
for (i in numnames) {
  data[,i] = as.numeric(data[,i])
}

data$dteday = as.Date(data$dteday)  #It changed date "02-04-11" to "2011-04-02".

str(data)

###_____Graphical analysis_____###
#Histogram for Target variable (continuous variable)
ggplot(data, aes_string(x = data$cnt)) +
  geom_histogram(fill="cornsilk", colour = "black") + geom_density() +
  scale_y_continuous(breaks=pretty_breaks(n=10)) +
  scale_x_continuous(breaks=pretty_breaks(n=10))+
  theme_bw() + xlab("Count of total rental bikes") + ylab("Frequency") + ggtitle("Target Variable Histogram") +
  theme(text=element_text(size=10))

#Histogram for Independent Continuous Variables
ggplot(data, aes_string(x = data$temp)) +
  geom_histogram(fill="blue", colour = "black") + geom_density() +
  scale_y_continuous(breaks=pretty_breaks(n=10)) +
  scale_x_continuous(breaks=pretty_breaks(n=10))+
  theme_bw() + xlab("Normalized Temperature (*C)") + ylab("Frequency") +
  ggtitle("IndependentVariable:Norm Temperature") +
  theme(text=element_text(size=10))
.
.
.
#And so on. The graphs are plotted and recorded in the project report.

###To extract days from "dteday" and make a new variable
data$day = day(data$dteday)
#As we already have information about the year and month, we have the whole date information & can
remove the "dteday" date type variable as it may not be suitable for modeling.
data[,1] = data[,16]
data[,16] = NULL      #dim = 731 x 15

```

```

col = names(data)

##### Missing Value Analysis #####
sum(is.na(data))

#There are no missing values for this data set.

##### Outlier Analysis #####
####Box Plot distribution & outlier check####
str(data)
for(i in 1:length(numnames)){
  assign(paste0("gn",i), ggplot(aes_string(y = (numnames[i]), x = data$cnt), data = subset(data))+
    stat_boxplot(geom = "errorbar", width = 0.5) +
    geom_boxplot(outlier.colour="red", fill = "light blue",outlier.shape=18,outlier.size=3, notch=FALSE) +
    theme(legend.position="bottom")+
    labs(y=numnames[i],x="Bike Rental Count")+
    ggtitle(paste("Box plot for",numnames[i])))
}

#Plotting plots together
gridExtra::grid.arrange(gn1,gn2,gn3,gn4,ncol=4)
gridExtra::grid.arrange(gn5,gn6,gn7,ncol=3)

#To check number of outliers in data (ignoring categorical variables, checked earlier)
out = 0.0
for(i in numnames){
  val = data[,i][data[,i] %in% boxplot.stats(data[,i])$out]
  out = out + length(val)
  print(i)
  print(length(val))
}
out #= 59. Total Outliers in the data set is 59.
#(59/731)*100 = 8.07% of data.

##To test for the best method to find missing values for this dataset
#data[12,12] #data[12,12] = 0.304627 (actual)
#data[12,12]= NA
#By median method:
#data$windspeed[is.na(data$windspeed)]=median(data$windspeed, na.rm = T)
#data[12,12] #data[12,12] = 0.180971 (median)

#reupload data
#data[12,12] #data[12,12] = 0.304627 (actual)
#data[12,12]= NA

```

```

#by mean method:
#data$windspeed[is.na(data$windspeed)]=mean(data$windspeed, na.rm = T)
#data[12,12] #data[12,12] = 0.1903299 (mean)

#reupload data
#data[12,12] #data[12,12] = 0.304627 (actual)
#data[12,12]= NA
#By KNN imputation method:
#(KNN takes only numeric inputs)
#for (i in col) {
# data[,i] = as.numeric(data[,i])
#}
#data= knnImputation(data, k=3) #For k=5,7,9, the difference was even more than k=3.
#data[12,12] #data[12,12] = 0.2324425 (KNN)
#We freeze NA imputation by MEDIAN method as it is closest to actual value.

#reupload data
#Converting outliers to NAs
#Select variables with outliers
Out_Var = c('hum','windspeed','casual') #Variables with outliers

for(i in Out_Var){
  val = data[,i][data[,i] %in% boxplot.stats(data[,i])$out]
  data[,i][data[,i] %in% val] = NA
}
sum(is.na(data)) #To verify

data= knnImputation(data, k=3)

sum(is.na(data)) #To verify

#Confirm again if any outlier exists
out = 0.0
for(i in numnames){
  val = data[,i][data[,i] %in% boxplot.stats(data[,i])$out]
  out= out + length(val)
  print(i)
  print(length(val))
}
out != 3. Windspeed has 2 outliers & Casual has 1 outlier.

#-->Redo 2 times the imputing using NAs by KNN imputation until 0 outliers.

write.csv(data, 'data_without Outliers.csv', row.names = F)

```

```

#To load the data
#data = read.csv("data_without Outliers.csv",header = T)

#####Feature Selection#####
#Correlation Plot
corrgram(data, order = F,
  upper.panel=panel.pie, text.panel=panel.txt, main = "Correlation Plot", font.labels = 1)
#cor(x), x must be numeric
#Convert all columns to numeric type
#for (i in col) {
# data[,i] = as.numeric(data[,i])
#} #NOTE: This changes all zero factor levels to numeric 1. so, "0" --> 1.
#mat = cor(data)
#corrplot(as.matrix(mat),method= 'pie',type = "lower", tl.col = "black", tl.cex = 0.7)
#If |r|>0.8, those two variables are redundant variables.
#Output: "mnth"-"season", "temp"-"atemp" & "cnt"-"registered" are highly positively correlated.
str(data)
#redo data conversion to proper types
catnames = c("season","yr","mnth","holiday","weekday","workingday","weathersit") #categorical variables
for (i in catnames) {
  data[,i] = as.factor(data[,i])
}

numnames = c("dteday","temp","atemp","hum","windspeed","casual","registered","cnt") #numerical
variables
for (i in numnames) {
  data[,i] = as.numeric(data[,i])
}

#####Chi-square Test of Independence (within Categorical Variables)
for(i in catnames){
  for(j in catnames){
    if(i!=j){
      print(names(data[i]))
      print(paste0(" Vs ", names(data[j])))
      print(chisq.test(table(data[,j],data[,i])))
    }
  }
}

#If p-value<0.05 (Reject Null Hypothesis) => variable A depends on variable B.
#If p-value>0.05 (Do Not Reject Null Hypothesis) => Variable A & variable B are independent of each other.
#Output: "workingday"-"holiday","weekday"-"workingday","weekday"-"holiday" & "mnth"-"season" depend
on each other significantly.

#####Using Random Forest Algorithm:

```

```

data.rf=randomForest(data$cnt~.,data = data, ntree=1000, keep.forest= F, importance= T)
importance(data.rf,type = 1)
#"holiday" has the least importance.
varImpPlot(data.rf,type = 1)

#####ANOVA test (comparision of Target Vs categorical variables)
anovacat = aov(cnt ~ season + yr + mnth + holiday + workingday + weekday + weathersit , data = data)
summary(anovacat)
#If p-value<0.05 (Reject Null Hypothesis) => Population means are significantly different.
#If p-value>0.05 (Do Not Reject Null Hypothesis) => Population means are not significantly different or are
same.

#####_____Feature Engineering_____#####
#From Chi-square test, we notice that "working day", "holiday" & "weekday" depend on each other and
intuitively there is a logical connection within them.
#We make a new variable using this connection between the three variables
#Denote: 1-->weekend, 2--> working day, 3--> holiday

data$day = NA
for (i in 1:nrow(data)){
  if ((data[i,7]=="0") && (data[i,5]=="0")){data[i,16] = 1}           #weekend
  else if ((data[i,7]=="1") && (data[i,5]=="0")){data[i,16] = 2}     #working day
  else if ((data[i,7]=="0") && (data[i,5]=="1")){data[i,16] = 3}     #holiday
  else data[i,16] =NA
}
sum(is.na(data$day)) #= 0, so no anomaly data case where it is working day & holiday both.

#####Dimensional Reduction#####
#Won't remove "dteday" variable as the user count is tracked on each day.
#As we added "day" new variable using "workingday" & "holiday", we can remove them both as "day" holds
the information of both.
data$holiday = data$day
data$day = NULL
colnames(data)[5] = "day"
data$day = as.factor(data$day)    # New variable "day": Factor w/ 3 levels "1","2","3"
#"Season" has multicollinearity problem as well and it is related to "mnth", so we can remove it.
data= subset(data, select= -c(season,workingday,temp,casual,registered))
factor_data = subset(data, select= c(yr,mnth,day,weekday,weathersit)) #5 factor variables
num_data = subset(data, select= c(dteday,atemp,hum,windspeed,cnt)) #5 numerical variables, contains
target variable
dim(data)          # 731 obs. x 10 variables
str(data)

#####Feature Scaling#####
#All continuous variables are already normalised in this data set.

```



```

rm(list= ls()[! (ls() %in% c('data','factor_data','num_data'))])

#####Sampling#####
set.seed(777)

sample.index = sample(nrow(data), 0.8*nrow(data), replace = F) #80% data -->Train set, 20%--> Test set
train = data[sample.index,]
test = data[-sample.index,]
dim(train) # 584 x 11
dim(test) # 147 x 11

#####Model Development#####
#As the target variable is of numeric type, this is a regression problem.
#####1.Decision Tree#####
#Decision trees can handle both categorical and numerical variables at the same time as features.
dt=rpart(cnt~.,data = train,method= "anova")
summary(dt)

#Predict for new test cases
predict.dt=predict(dt,test[,-10])

#Error metric:
postResample(predict.dt,test[,10])
#Output:
#RMSE    Rsquared    MAE
#1036.8218286  0.7105788  768.8217306

#calculate MAPE
mape = function(y,yi)
{mean(abs((y-yi)/y))*100
}
mape.dt = mape(test[,10],predict.dt) #30.79%

library(mltools)
rmsle(predict.dt,test[,10]) #0.3665

#####2.Random Forest Algorithm#####
rf = randomForest(cnt~., train, importance = TRUE, ntree = 500)
summary(rf)
#Predict for test case:
predict.rf <- data.frame(predict(rf, subset(test, select = -c(cnt))))
#Error metric:
postResample(predict.rf,test[,10])
#Output:

```

```

#RMSE      Rsquared    MAE
#778.4675527  0.8507608 576.6110231

mape.rf = mape(test[,10],predict.rf$predict.rf..subset.test..select....c.cnt...) # 24.9%

#####3.Multiple Linear Regression#####

#creating dummy variables for categorical data
library(dummies)
factor_new = dummy.data.frame(factor_data, sep = ".") #731 x 27

#sampling#
df = cbind(factor_new, num_data)
#for (i in 1:ncol(df)) {
#  df[,i] = as.numeric(df[,i])
#}
str(df)      # 731 X 32

set.seed(123)
train_index = sample(1:nrow(df), 0.8*nrow(df))
train.df = df[train_index,]      #584 x 32
test.df = df[-train_index,]      #147 x 32

#Check Multicollinearity
library(usdm)
vif(df[, -32])
vifcor(df[, -32], th = 0.8)
#Output:
#3 variables from the 31 input variables have collinearity problem: yr.1, weathersit.2, day.2
#removing multicollinear variables and redo check:
df = subset(df, select= -c(yr.1, weathersit.2, day.2))
train.df = subset(train.df, select= -c(yr.1, weathersit.2, day.2)) #584 x 29
test.df = subset(test.df, select= -c(yr.1, weathersit.2, day.2)) #147 x 29
dim(df) #731 x 29
#Recheck VIFCORR: No variable from the 29 input variables has collinearity problem.

#run regression model
lr = lm(cnt~., data = train.df)
#summary of the model
summary(lr)

#Predict for test case:
predict.lr= predict(lr, test.df[, -29])
#Error metric:
postResample(predict.lr, test.df[, 29])

```

```

#Output:
#RMSE      Rsquared    MAE
#800.2783046  0.8303233 581.4298996

mape.lr = mape(test.df[,29],predict.lr)    #17.5%

#####4.KNN Implementation#####
#To check for best k value:
model <- train(cnt~., data = train, method = "knn",
               trControl = trainControl("cv", number = 10),
               tuneLength = 15)
model$bestTune
#k = 2 , 7
plot(model)

#K=3:
predict.knn = knn.reg(train = train.df[, -29], test = test.df[, -29], train.df$cnt, k = 3)
print(predict.knn)

#Error metric:
postResample(predict.knn$pred, test.df[, 29])
#Output:
#RMSE      Rsquared    MAE
#1392.7631351  0.4544424 1110.0045351

mape.knn = mape(test.df[, 29], predict.knn$pred) #38.98%

#K=5:
#predict.knn = knn.reg(train = train.df[, -29], test = test.df[, -29], train.df$cnt, k = 5)
#print(predict.knn)
#Error metric:
#mape(test.df[, 29], predict.knn$pred)
#Output:
#mape
#45.26592 %
#postResample(predict.knn$pred, test.df[, 29])
#RMSE      Rsquared    MAE
#1450.9419952  0.4484269 1169.3782313

#K=7:
#predict.knn = knn.reg(train = train.df[, -29], test = test.df[, -29], train.df$cnt, k = 7)
#print(predict.knn)
#Error metric:
#mape(test.df[, 29], predict.knn$pred)
#Output:

```

```
#mape
#47.63637 %
#postResample(predict.knn$pred,test.df[,29])
#RMSE      Rsquared      MAE
#1456.0507716  0.4983456 1171.8736638
#####And so on, done upto k = 11.

#A new dataframe to store results
algorithm <- c('Decision Tree','Random Forest','Linear Regression','KNN')
MAPE_val <- c(mape.dt,mape.rf,mape.lr,mape.knn)
results <- data.frame(algorithm, MAPE_val)
print(results)
barplot(results$MAPE_val, width = 1, names.arg = results$algorithm,
        ylab="MAPE value", xlab = "Algorithm",col='pink')

##Thus, we find the "Multiple Linear Regression Algorithm" gives us the best result with the least MAPE for
this dataset.
```

Appendix C: Python code

```
#Set working directory
import os
os.chdir("F:/DS/edWisor/Project 2")
os.getcwd()
```

Load libraries

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
from random import randrange, uniform
from scipy.stats import chi2_contingency
from ggplot import *
```

In []:

```
from fancyimpute import KNN
```

In []:

```
import datetime as dt
```

In []:

```
#Load the data
data = pd.read_csv("day.csv")
```

In []:

Data exploration

In []:

```

data.shape
In [ ]:

data.head(10)
In [ ]:

type(data)
In [ ]:

data.info()
In [ ]:

#Missing Value Analysis
#Check for missing value
data.isnull().sum()
#No missing values in the dataset
In [ ]:

#remove "instant" variable as it is just like serial number & doesn't predict
data = data.drop(['instant'], axis=1)
In [ ]:

data.shape
In [ ]:

#extracting day number from 'dteday' variable
data['dteday'].apply(str)
data['dteday'] = pd.to_datetime(data['dteday'])
data['dteday'] = pd.DatetimeIndex(data['dteday']).day
#removing 'dteday' variable
In [ ]:

data.head(20)
In [ ]:

#save numeric & categorical names
numnames = ["dteday", "temp", "atemp", "hum", "windspeed", "casual", "registered", "cnt"]
catnames = ["season", "yr", "mnth", "holiday", "weekday", "workingday", "weathersit"]
data.shape
In [ ]:

for i in catnames:
    data[i] = data[i].astype('object')
for i in numnames:
    data[i] = data[i].astype('float')
In [ ]:

data.dtypes
In [ ]:



## Outlier analysis


In [ ]:

#Plot boxplot to visualize Outliers
%matplotlib inline
plt.boxplot(data['windspeed'])
In [ ]:

#Detect and delete outliers from data

```

```

for i in numnames:
    print(i)
    q75, q25 = np.percentile(data.loc[:,i], [75 ,25])
    iqr = q75 - q25

    min = q25 - (iqr*1.5)
    max = q75 + (iqr*1.5)
    print(min)
    print(max)

    #Remove the outliers
    data = data.drop(data[data.loc[:,i] < min].index)
    data = data.drop(data[data.loc[:,i] > max].index)

    #data.loc[data[i] < min,:i] = np.nan
    #data.loc[data[i] > max,:i] = np.nan

#Calculate missing value
#missing_val = pd.DataFrame(data.isnull().sum())
#Impute with KNN
#data = pd.DataFrame(KNN(21).fit_transform(data), columns = data.columns)

```

```
data.shape      #55 rows deleted
```

In []:

```
data.isnull().sum()
```

In []:

Feature Selection

```

##Correlation analysis
#Correlation plot
df_corr = data.loc[:,numnames]
#Set the width and height of the plot
f, ax = plt.subplots(figsize=(7, 5))

#Generate correlation matrix
corr = df_corr.corr()

#Plot using seaborn library
sns.heatmap(corr, mask=np.zeros_like(corr, dtype=np.bool), cmap=sns.diverging_palette(220, 10, as_cmap=True),
            square=True, ax=ax)

```

```

#Chisquare test of independence
#loop for chi square values
for i in catnames:
    print(i)
    chi2, p, dof, ex = chi2_contingency(pd.crosstab(data['cnt'], data[i]))
    print(p)

```

In []:

```

In [ ]:
#New Categorical Variable containing the data of "workingday" & "holiday"
#Denote: 1-->weekend, 2--> working day, 3--> holiday
data.loc[(data['workingday'] == 0) & (data['holiday'] == 0), 'day'] = '1'
data.loc[(data['workingday'] == 1) & (data['holiday'] == 0), 'day'] = '2'
data.loc[(data['workingday'] == 0) & (data['holiday'] == 1), 'day'] = '3'

In [ ]:
data = data.drop(["workingday", "holiday", "temp", "casual", "registered"], axis=1)

In [ ]:
data.head(10)

In [ ]:
df = data[['dteday', 'mnth', 'yr', 'season', 'weekday', 'day', 'weathersit', 'atemp', 'hum', 'windspeed', 'cnt']]

In [ ]:
df.head(10)

In [ ]:
#####Feature Scaling#####
#All continuous variables are already normalised in this data set.

numnames = ["dteday", "atemp", "hum", "windspeed"]    #not including "cnt" target variable
catnames = ["mnth", "yr", "season", "weekday", "day", "weathersit"]



## Model Development



In [ ]:
#Data Sampling
nrow= len(df.index)
train, test = train_test_split(df, test_size = 0.2)

In [ ]:
train.shape    #540 x 11
test.shape     #136 x 11

In [ ]:
#####Decision Tree Algorithm
from sklearn.tree import DecisionTreeRegressor
fit_dt= DecisionTreeRegressor(max_depth=2).fit(train.iloc[:,0:10],train.iloc[:,10])

In [ ]:
fit_dt

In [ ]:
predict_dt= fit_dt.predict(test.iloc[:,0:10])

In [ ]:
#Calculate RMSE
def RMSE(actual, pred):
    return np.sqrt(((pred - actual) ** 2).mean())

RMSE(test.iloc[:,10],predict_dt)
#output = 1162.84440171958

```

```

#####Random Forest Algorithm
from sklearn.ensemble import RandomForestRegressor
fit_rf = RandomForestRegressor(n_estimators = 100, random_state = 99).fit(train.iloc[:,0:10],train.iloc[:,10])

fit_rf

predict_rf= fit_rf.predict(test.iloc[:,0:10])

RMSE(test.iloc[:,10],predict_rf)
#output = 765.0407919968172

#####Multiple Linear Regression
import statsmodels.api as sm
#Creat dataframe with all numerical variables
df.lr = df[['cnt','dteday','atemp','hum','windspeed']]
#create dummies for categorical variables
for i in catnames:
    temp = pd.get_dummies(df[i],prefix = i)
    df.lr = df.lr.join(temp)

df.lr.shape          #676 x 36

#split data into train-test sets
s = np.random.rand(len(df.lr))<0.8
train.lr = df.lr[s]    #80%
test.lr = df.lr[~s]    #20%

train.lr.shape        #564 x 36
test.lr.shape         #112 x 36

#Build MLR model
fit_lr = sm.OLS(train.lr.iloc[:,0],train.lr.iloc[:,1:35]).fit()
fit_lr.summary()

predict_lr = fit_lr.predict(test.lr.iloc[:,1:35])

RMSE(test.lr.iloc[:,0],predict_lr)
#output = 713.1957640471251

#####KNN Implementation
from sklearn import neighbors
rmse_val = []          #to store rmse values for different k
for K in range(30):
    K = K+1

```



```

fit_knn = neighbors.KNeighborsRegressor(n_neighbors = K)

fit_knn.fit(train.iloc[:,0:10], train.iloc[:,10]) #fit the model

predict_knn = fit_knn.predict(test.iloc[:,0:10]) #make prediction on test set
error = RMSE(test.iloc[:,10] , predict_knn) #calculate rmse
rmse_val.append(error) #store rmse values
print('RMSE value for k= ', K , 'is:', error)

```

In []:

#plotting the rmse values against k values

```

curve = pd.DataFrame(rmse_val)
curve.plot()

```

#K=2 is the value of neighbors for least RMSE.

In []:

#For K=12:

```

fit_knn = neighbors.KNeighborsRegressor(n_neighbors = 2)
fit_knn.fit(train.iloc[:,0:10], train.iloc[:,10]) #fit the model
predict_knn = fit_knn.predict(test.iloc[:,0:10]) #make prediction on test set
RMSE(test.iloc[:,10] , predict_knn)

```

#output = 1209.595772142617

In []:

#Thus, we find the "Multiple Linear Regression Algorithm" gives us the best result with the least RMSE for this dataset.

References

Mitchell, T. (1997). Machine Learning. McGraw Hill. p. 2

Brieman, Friedman, Olshen and Stone, Classification and Regression Trees, 1984